

Web of Things (WoT) Architecture

W3C Editor's Draft 18 June 2020



This version:

<https://w3c.github.io/wot-architecture/>

Latest published version:

<https://www.w3.org/TR/wot-architecture/>

Latest editor's draft:

<https://w3c.github.io/wot-architecture/>

Editors:

Matthias Kovatsch ([Huawei](#))

Ryuichi Matsukura ([Fujitsu Ltd.](#))

Michael Lagally ([Oracle Corp.](#))

Toru Kawaguchi ([Panasonic Corp.](#))

Kunihiko Toumura ([Hitachi, Ltd.](#))

Kazuo Kajimoto (Former Editor, when at Panasonic)

Participate:

[GitHub w3c/wot-architecture](#)

[File a bug](#)

[Commit history](#)

[Pull requests](#)

Contributors:

[In the GitHub repository](#)

Copyright © 2017-2020 [W3C](#)[®] ([MIT](#), [ERCIM](#), [Keio](#), [Beihang](#)). W3C [liability](#), [trademark](#) and [permissive document license](#) rules apply.

Abstract

The W3C Web of Things (WoT) is intended to enable interoperability across IoT platforms and application domains. Overall, the goal of the WoT is to preserve and complement existing IoT standards and solutions. In general, the W3C WoT architecture is designed to describe what exists rather than to prescribe what to implement.

This *WoT Architecture* specification describes the abstract architecture for the [W3C](#) Web of Things. This abstract architecture is based on a set of requirements that were derived from use cases for multiple application domains, both given in this document. A set of modular building blocks are also identified whose detailed specifications are given in other documents. This document describes how these building blocks are related and work together. The WoT abstract architecture defines a basic conceptual framework that can be mapped onto a variety of concrete deployment scenarios, several examples of which are given. However, the abstract architecture described in this specification does not itself define concrete mechanisms or prescribe any concrete implementation.

Status of This Document

This section describes the status of this document at the time of its publication. Other documents may supersede this document. A list of current [W3C](#) publications and the latest revision of this technical report can be found in the [W3C technical reports index](#) at <https://www.w3.org/TR/>.

This document describes an abstract architecture design. However, there is an [Implementation Report](#) that describes a set of concrete implementations based on the associated *WoT Thing Description* specification. These are implementations following the [W3C](#) Web of Things architecture.

This document was published by the [Web of Things Working Group](#) as an Editor's Draft.

[GitHub Issues](#) are preferred for discussion of this specification. Alternatively, you can send comments to our mailing list. Please send them to public-wot-wg@w3.org ([archives](#)).

Publication as an Editor's Draft does not imply endorsement by the [W3C](#) Membership. This is a draft document and may be updated, replaced or obsoleted by other documents at any time. It is inappropriate to cite this document as other than work in progress.

This document was produced by a group operating under the [W3C Patent Policy](#). [W3C](#) maintains a [public list of any patent disclosures](#) made in connection with the deliverables of the group; that page also includes instructions for disclosing a patent. An individual who has actual knowledge of a patent which the individual believes contains [Essential Claim\(s\)](#) must disclose the information in accordance with [section 6 of the W3C Patent Policy](#).

This document is governed by the [1 March 2019 W3C Process Document](#).

Table of Contents

1. Introduction

2. Conformance

3. Terminology

4. Use Cases

4.1 Application Domains

- 4.1.1 Consumer
- 4.1.2 Industrial
 - 4.1.2.1 Example: Smart Factory
- 4.1.3 Transportation & Logistics
- 4.1.4 Utilities
- 4.1.5 Oil and Gas
- 4.1.6 Insurance
- 4.1.7 Engineering and Construction
- 4.1.8 Agriculture
- 4.1.9 Healthcare
- 4.1.10 Environment Monitoring
- 4.1.11 Smart Cities
- 4.1.12 Smart Buildings
- 4.1.13 Connected Car
 - 4.1.13.1 Connected Car Example

4.2 Common Patterns

- 4.2.1 Device Controllers
- 4.2.2 Thing-to-Thing
- 4.2.3 Remote Access
- 4.2.4 Smart Home Gateways
- 4.2.5 Edge Devices
- 4.2.6 Digital Twins
 - 4.2.6.1 Cloud-ready Devices
 - 4.2.6.2 Legacy Devices
- 4.2.7 Multi-Cloud
- 4.2.8 Cross-domain Collaboration

4.3 Summary

5. Requirements

5.1 Functional Requirements

- 5.1.1 Common Principles
- 5.1.2 Thing Functionalities
- 5.1.3 Search and Discovery

- 5.1.4 Description Mechanism
- 5.1.5 Description of Attributes
- 5.1.6 Description of Functionalities
- 5.1.7 Network
- 5.1.8 Deployment
- 5.1.9 Application
- 5.1.10 Legacy Adoption
- 5.2 Technical Requirements
 - 5.2.1 Components in the Web of Things and the Web of Things Architecture
 - 5.2.2 Devices
 - 5.2.3 Applications
 - 5.2.4 Digital Twins
 - 5.2.5 Discovery
 - 5.2.6 Security
 - 5.2.7 Accessibility

6. WoT Architecture

- 6.1 Overview
- 6.2 Affordances
- 6.3 Web Thing
- 6.4 Interaction Model
 - 6.4.1 Properties
 - 6.4.2 Actions
 - 6.4.3 Events
- 6.5 Hypermedia Controls
 - 6.5.1 Links
 - 6.5.2 Forms
- 6.6 Protocol Bindings
 - 6.6.1 Hypermedia-driven
 - 6.6.2 URIs
 - 6.6.3 Standard Set of Methods
 - 6.6.4 Media Types
- 6.7 WoT System Components and Their Interconnectivity
 - 6.7.1 Direct Communication
 - 6.7.2 Indirect Communication

7. WoT Building Blocks

- 7.1 WoT Thing Description
- 7.2 WoT Binding Templates

7.3	WoT Scripting API
7.4	WoT Security and Privacy Guidelines
8.	Abstract Servient Architecture
8.1	Behavior Implementation
8.2	WoT Runtime
8.3	WoT Scripting API
8.4	Exposed Thing and Consumed Thing Abstractions
8.5	Private Security Data
8.6	Protocol Stack Implementation
8.7	System API
8.8	Alternative Servient and WoT Implementations
8.8.1	Native WoT API
8.8.2	Thing Description for Existing Devices
9.	Example WoT Deployments
9.1	Thing and Consumer Roles
9.2	Topology of WoT Systems and Deployment Scenarios
9.2.1	Consumer and Thing on the Same Network
9.2.2	Consumer and Thing Connected via Intermediaries
9.2.2.1	Intermediary Acting as a Proxy
9.2.2.2	Intermediary Acting as a Digital Twin
9.2.3	Devices in a Local Network Controlled from a Cloud Service
9.2.4	Discovery Using a Thing Directory
9.2.5	Service-to-Service Connections Across Multiple Domains
9.2.5.1	Connection Through Thing Directory Synchronization
9.2.5.2	Connection Through Proxy Synchronization
10.	Security and Privacy Considerations
10.1	WoT Thing Description Risks
10.1.1	Thing Description Private Security Data Risk
10.1.2	Thing Description Personally Identifiable Information Risk
10.1.3	Thing Description Communication Metadata Risk
10.2	WoT Scripting API Security and Privacy Risks
10.2.1	Cross-Script Security and Privacy Risk
10.2.2	Physical Device Direct Access Security and Privacy Risk
10.3	WoT Runtime Security and Privacy Risks
10.3.1	Provisioning and Update Security Risk
10.3.2	Security Credentials Storage Security and Privacy Risk

A. Recent Specification Changes

B. Acknowledgments

C. References

C.1 Normative references

C.2 Informative references

1. Introduction §

The goals of the *Web of Things* (WoT) are to improve the interoperability and usability of the Internet of Things (IoT). Through a collaboration involving many stakeholders over many years, several building blocks have been identified that help address these challenges.

This specification is focused on the scope of W3C WoT standardization, which can be broken down into these building blocks as well as the abstract architecture that defines how they are related. The building blocks are defined and described in detail in separate specifications. However, in addition to defining the abstract architecture and its terminology and conceptual framework, this specification also serves as an introduction to the WoT building blocks, and explains their interworking:

- The *Web of Things (WoT) Thing Description* [[WOT-THING-DESCRIPTION](#)] normatively provides a machine-readable data format for describing the metadata and network-facing interfaces of Things. It is based upon the fundamental concepts introduced in this document, such as interaction affordances.
- The *Web of Things (WoT) Binding Templates* [[WOT-BINDING-TEMPLATES](#)] provides informational guidelines on how to define network-facing interfaces in Things for particular protocols and IoT ecosystems, which we call Protocol Bindings. The document also provides examples for a number of existing IoT ecosystems and standards.
- The *Web of Things (WoT) Scripting API* [[WOT-SCRIPTING-API](#)], which is optional, enables the implementation of the application logic of a Thing using a common JavaScript API similar to the Web browser APIs. This simplifies IoT application development and enables portability across vendors and devices.
- The *Web of Things (WoT) Security and Privacy Guidelines* [[WOT-SECURITY](#)] represent a cross-cutting building block. This informational document provides guidelines for the secure implementation and configuration of Things, and discusses issues which should be considered in any systems implementing W3C WoT. However, it should be emphasized that security and privacy can only be fully evaluated in the context of a complete set of concrete mechanisms for a specific implementation, which the WoT abstract architecture does not fully specify. This is

especially true when the WoT architecture is used descriptively for pre-existing systems, since the W3C WoT cannot constrain the behavior of such systems, it can only describe them. In this document we also discuss privacy and security risks and their mitigation at a high level in section [§ 10. Security and Privacy Considerations](#).

This specification also covers non-normative architectural aspects and conditions for the deployment of WoT systems. These guidelines are described in the context of example deployment scenarios, although this specification does not normatively define specific concrete implementations.

This specification serves as an umbrella for W3C WoT specifications and defines the basics such as terminology and the underlying abstract architecture of the W3C Web of Things. In summary, the purpose of this specification is to provide:

- a set of use cases in [§ 4. Use Cases](#) that lead to the W3C WoT Architecture,
- a set of requirements for WoT implementations in [§ 5. Requirements](#),
- a definition of the abstract architecture in [§ 6. WoT Architecture](#)
- an overview of a set of WoT building blocks and their interplay in [§ 7. WoT Building Blocks](#),
- an informative guideline on how to map the abstract architecture to possible concrete implementations in [§ 8. Abstract Servient Architecture](#),
- informative examples of possible deployment scenarios in [§ 9. Example WoT Deployments](#),
- and a discussion, at a high level, of security and privacy considerations to be aware of when implementing a system based on the W3C WoT architecture in [§ 10. Security and Privacy Considerations](#).

Additional requirements, use cases, conceptual features and new building blocks will be addressed in a future revision of this document.

2. Conformance §

As well as sections marked as non-normative, all authoring guidelines, diagrams, examples, and notes in this specification are non-normative. Everything else in this specification is normative.

The key words *MAY*, *MUST*, and *SHOULD* in this document are to be interpreted as described in [BCP 14 \[RFC2119\]](#) [\[RFC8174\]](#) when, and only when, they appear in all capitals, as shown here.

3. Terminology §

This section is non-normative.

This specification uses the following terms as defined here. The WoT prefix is used to avoid ambiguity for terms that are (re)defined specifically for Web of Things concepts.

Action

An Interaction Affordance that allows to invoke a function of the Thing, which manipulates state (e.g., toggling a lamp on or off) or triggers a process on the Thing (e.g., dim a lamp over time).

Binding Templates

A re-usable collection of blueprints for the communication with different IoT platforms. The blueprints provide information to map Interaction Affordances to platform-specific messages through WoT Thing Description as well as implementation notes for the required protocol stacks or dedicated communication drivers.

Consumed Thing

A software abstraction that represents a remote Thing used by the local application. The abstraction might be created by a native WoT Runtime, or instantiated as an object through the WoT Scripting API.

Consuming a Thing

To parse and process a TD document and from it create a Consumed Thing software abstraction as interface for the application in the local runtime environment.

Consumer

An entity that can process WoT Thing Descriptions (including its JSON-based representation format) and interact with Things (i.e., consume Things).

Data Schema

A data schema describes the information model and the related payload structure and corresponding data items that are passed between Things and Consumers during interactions.

Digital Twin

A digital twin is a virtual representation of a device or a group of devices that resides on a cloud or edge node. It can be used to represent real-world devices which may not be continuously online, or to run simulations of new applications and services, before they get deployed to the real devices.

Domain-specific Vocabulary

Linked Data vocabulary that can be used in the WoT Thing Description, but is not defined by W3C WoT.

Edge Device

A device that provides an entry point into enterprise or service provider core networks. Examples include gateways, routers, switches, multiplexers, and a variety of other access devices.

Event

An Interaction Affordance that describes an event source, which asynchronously pushes event data to Consumers (e.g., overheating alerts).

Exposed Thing

A software abstraction that represents a locally hosted Thing that can be accessed over the network by remote Consumers. The abstraction might be created by a native WoT Runtime, or instantiated as an object through the WoT Scripting API.

Exposing a Thing

To create an Exposed Thing software abstraction in the local runtime environment to manage the state of a Thing and interface with the behavior implementation.

Hypermedia Control

A serialization of a Protocol Binding in hypermedia, that is, either a Web link [[RFC8288](#)] for navigation or a Web form for performing other operations. Forms can be seen as request templates provided by the Thing to be completed and sent by the Consumer.

Interaction Affordance

Metadata of a Thing that shows and describes the possible choices to Consumers, thereby suggesting how Consumers may interact with the Thing. There are many types of potential affordances, but ~~W3C~~ WoT defines three types of Interaction Affordances: Properties, Actions, and Events. A fourth Interaction Affordance is navigation, which is already available on the Web through linking.

Interaction Model

An intermediate abstraction that formalizes and narrows the mapping from application intent to concrete protocol operations. In ~~W3C~~ WoT, the defined set of Interaction Affordances constitutes the Interaction Model.

Intermediary

An entity between Consumers and Things that can proxy, augment, or compose Things and republish a WoT Thing Description that points to the WoT Interface on the Intermediary instead of the original Thing. For Consumers, an Intermediary may be indistinguishable from a Thing, following the Layered System constraint of REST.

IoT Platform

A specific IoT ecosystem such as OCF, oneM2M, or Mozilla Project Things with its own specifications for application-facing APIs, data model, and protocols or protocol configurations.

Metadata

Data that provides a description of an entity's abstract characteristics. For example, a [Thing Description](#) is Metadata for a [Thing](#).

Personally Identifiable Information (PII)

Any information that can be used to identify the natural person to whom such information relates, or is or might be directly or indirectly linked to a natural person. We use the same definition as [[ISO-IEC-29100](#)].

Privacy

Freedom from intrusion into the private life or affairs of an individual when that intrusion results from undue or illegal gathering and use of data about that individual. We use the same definition as [\[ISO-IEC-2382\]](#). See also [Personally Identifiable Information](#) and [Security](#), as well as other related definitions in [\[ISO-IEC-29100\]](#).

Private Security Data

Private Security Data is that component of a Thing's Security Configuration that is kept secret and is not shared with other devices or users. An example would be private keys in a PKI system. Ideally such data is stored in a separate memory inaccessible to the application and is only used via abstract operations, such as signing, that do not reveal the secret information even to the application using it.

Property

An Interaction Affordance that exposes state of the Thing. This state can then be retrieved (read) and optionally updated (write). Things can also choose to make Properties observable by pushing the new state after a change.

Protocol Binding

The mapping from an Interaction Affordance to concrete messages of a specific protocol, thereby informing Consumers how to activate the Interaction Affordance. [W3C WoT](#) serializes Protocol Bindings as hypermedia controls.

Public Security Metadata

Public Security Metadata is that component of a Thing's Security Configuration which describes the security mechanisms and access rights necessary to access a Thing. It does not include any secret information or concrete data (including public keys), and does not by itself, provide access to the Thing. Instead, it describes the mechanisms by which access may be obtained by authorized users, including how they must authenticate themselves.

Security

Preservation of the confidentiality, integrity and availability of information. Properties such as authenticity, accountability, non-repudiation, and reliability may also be involved. This definition is adapted from the definition of *Information Security* in [\[ISO-IEC-27000\]](#), which also includes additional definitions of each of the more specific properties mentioned. Please refer to this document for other related definitions. We additionally note that it is desirable that these properties be maintained both in normal operation and when the system is subject to attack.

Security Configuration

The combination of Public Security Metadata, Private Security Data, and any other configuration information (such as public keys) necessary to operationally configure the security mechanisms of a Thing.

Servient

A software stack that implements the WoT building blocks. A Servient can host and expose Things and/or host Consumers that consume Things. Servients can support multiple Protocol Bindings to enable interaction with different IoT platforms.

Subprotocol

An extension mechanism to a transfer protocol that must be known to interact successfully. An example is long polling for HTTP.

TD

Short for WoT Thing Description.

TD Vocabulary

A controlled Linked Data vocabulary by W3C WoT to tag the metadata of Things in the WoT Thing Description including communication metadata of WoT Binding Templates.

Thing or Web Thing

An abstraction of a physical or a virtual entity whose metadata and interfaces are described by a WoT Thing Description, whereas a virtual entity is the composition of one or more Things.

Thing Directory

A directory service for TDs that provides a Web interface to register TDs (similar to [\[CoRE-RD\]](#)) and look them up (e.g., using SPARQL queries or the CoRE RD lookup interface [\[CoRE-RD\]](#)).

Transfer Protocol

The underlying, standardized application layer protocol without application-specific requirements or constraints on options or subprotocol mechanisms. Examples are HTTP, CoAP, or MQTT.

Virtual Thing

An instance of a Thing that represents a Thing that is located on another system component.

WoT Interface

The network-facing interface of a Thing that is described by a WoT Thing Description.

WoT Runtime

A runtime system that maintains an execution environment for applications, and is able to expose and/or consume Things, to process WoT Thing Descriptions, to maintain Security Configurations, and to interface with Protocol Binding implementations. A WoT Runtime may have a custom API or use the optional WoT Scripting API.

WoT Scripting API

The application-facing programming interface provided by a Servient in order to ease the implementation of behavior or applications running in a WoT Runtime. It is comparable to the Web browser APIs. The WoT Scripting API is an optional building block for W3C WoT.

WoT Servient

Synonym for Servient.

WoT Thing Description or Thing Description

Structured data describing a Thing. A WoT Thing Description comprises general metadata, domain-specific metadata, Interaction Affordances (which include the supported Protocol Bindings), and links to related Things. The WoT Thing Description format is the central building block of W3C WoT.

4. Use Cases §

This section is non-normative.

This section presents the application domains and use cases targeted by the W3C WoT and which are used to derive the abstract architecture discussed in [§ 7. WoT Building Blocks](#).

The Web of Things architecture does not put any limitations on use cases and application domains. Various application domains have been considered to collect common patterns that have to be satisfied by the abstract architecture.

The following sections are not exhaustive. Rather they serve as illustrations, where connected things can provide additional benefit or enable new scenarios.

4.1 Application Domains §

4.1.1 Consumer §

In the consumer space there are multiple assets that benefit from being connected. Lights and air conditioners can be turned off based on room occupancy. Window blinds can be closed automatically based on weather conditions and presence. Energy and other resource consumption can be optimized based on usage patterns and predictions.

The consumer use cases in this section includes Smart Home use case.

[Figure 1](#) shows an example of a Smart Home. In this case, gateways are connected to edge devices such as sensors, cameras and home appliances through corresponding local communication protocols such as KNX, ECHONET, ZigBee, DECT ULE and Wi-SUN. Multiple gateways can exist in one home, while each gateway can support multiple local protocols.

Gateways can be connected to the cloud through the internet, while some appliances can be connected to the cloud directly. Services running in the cloud collect data from edge devices and analyze the data, then provide value to users through the edge devices and other UX devices.

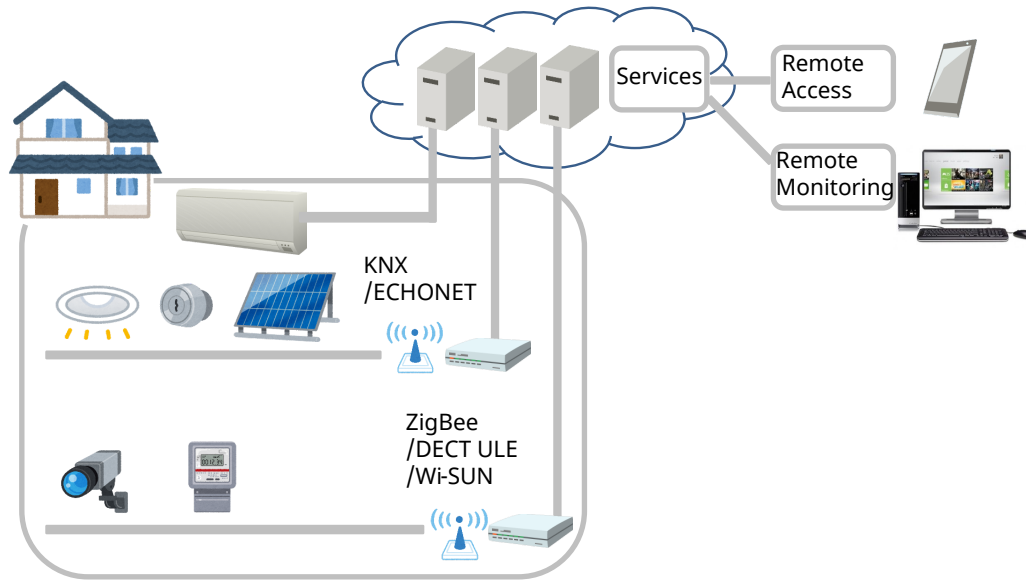


Figure 1 Smart Home

Smart home provides consumer benefits such as remote access and control, voice control and home automation. Smart home also enables device manufacturers to monitor and maintain devices remotely. Smart home can realize value added services such as energy management and security surveillance.

4.1.2 Industrial §

The industrial use cases in this section are applicable to different industry verticals. Due to the nature of overlaps in the application scenarios, different verticals have similar use cases.

4.1.2.1 Example: Smart Factory §

[Figure 2](#) shows an example of a Smart Factory. In this case, field-level, cell and line controllers automate different factory equipment based on industrial communication protocols such as PROFINET, Modbus, OPC UA TSN, EtherCAT, or CAN. An industrial edge device collects selected data from various controllers and makes it available to a cloud backend service, e.g., for remote monitoring via a dashboard or analyzes it for preventive maintenance.




Figure 2 Smart Factory

Smart factories require advanced monitoring of the connected manufacturing equipment as well of the manufactured products. They benefit from predictions of machine failures and early discovery of anomalies to prevent costly downtime and maintenance efforts.

Additionally, monitoring of connected manufacturing equipment and the environment at the production facility for the presence of poisonous gases, excessive noise or heat increases the safety of the workers and reduces the risks of incidents or accidents.

Real-time monitoring and KPI calculations of production equipment helps to detect productivity problems and optimize the supply chain.

4.1.3 Transportation & Logistics §

Monitoring of vehicles, fuel costs, maintenance needs and assignments helps to optimize the full utilization of the vehicle fleet.

Shipments can be tracked to be en-route to ensure consistent quality and condition of the transported goods. This is especially useful to assert the integrity of the cold-chain from warehouses to refrigerated trucks to delivery.

Centralized monitoring and management of stock in warehouses and yards can prevent out of stock and excessive stock situations.

4.1.4 Utilities §

Automated reading of residential and C&I (Commercial and Industrial) meters, and billing offers continuous insights into resource consumption and potential bottlenecks.

Monitoring the condition and output of distributed renewable energy generation equipment enables optimization of distributed energy resources.

Monitoring and remote-controlling of distribution equipment helps to automate the distribution process.

Continuous monitoring of generation and distribution infrastructure is improving safety of utilities crew in the field.

4.1.5 Oil and Gas §

Offshore platform monitoring, leakage detection and prediction of pipelines as well as monitoring and controlling the levels in tanks and reservoirs helps to improve the industrial safety for the workforce as well as for the environment.

Automated calculation of a distributed stock through various storage tanks and delivery pipes/trucks allows for improved planning and resource optimization.

4.1.6 Insurance §

Proactive Asset Monitoring of high value assets such as connected structures, fleet vehicles, etc. mitigates the risk of severe damage and high costs due to predictions and early detection of incidents.

Usage based insurance can be offered with usage tracking and customized insurance policies.

Predictive weather monitoring and re-routing fleet vehicles to covered garages can limit loss due to hail damage, tree damage.

4.1.7 Engineering and Construction §

Monitoring for industrial safety reduces the risks of security hazards. Monitoring of assets at construction site can prevent damage and loss.

4.1.8 Agriculture §

Soil condition monitoring and creating optimal plans for watering, fertilizing as well as monitoring the produce conditions optimize the quality and output of agricultural produce.

4.1.9 Healthcare §

Data collection and analytics of clinical trial data helps to gain insights into new areas.

Remote patient monitoring mitigates the risk of undetected critical situations for elderly people and patients after hospitalization.

4.1.10 Environment Monitoring §

Environment monitoring typically relies on a lot of distributed sensors that send their measurement data to common gateways, edge devices and cloud services.

Monitoring of air pollution, water pollution and other environmental risk factors such as fine dust, ozone, volatile organic compound, radioactivity, temperature, humidity to detect critical environment conditions can prevent unrecoverable health or environment damages.

4.1.11 Smart Cities §

Monitoring of Bridges, Dams, Levees, Canals for material condition, deterioration, vibrations discovers maintenance repair work and prevents significant damage. Monitoring of highways and providing appropriate signage ensures optimized traffic flow.

Smart Parking is optimizing and tracking the usage and availability of parking spaces and automates billing/reservations.

Smart control of street lights based on presence detection, weather predictions, etc. reduces cost.

Garbage containers can be monitored to optimize the waste management and the trash collection route.

4.1.12 Smart Buildings §

Monitoring the energy usage throughout the building helps to optimize resource consumption and reduce waste.

Monitoring the equipment in the buildings such as HVAC, Elevators, etc. and fixing problems early improves the satisfaction of occupants.

4.1.13 Connected Car §

Monitoring of operation status, prediction of service needs optimizes maintenance needs and costs. Driver safety is enhanced with notifications of an early warning system for critical road and traffic conditions.

4.1.13.1 Connected Car Example §

[Figure 3](#) shows an example of a Connected Car. In this case, a gateway connects to car components through CAN and to the car navigation system through a proprietary interface. Services running in the cloud collect data pushed from car components and analyze the data from multiple cars to determine traffic patterns. The gateway can also consume cloud services, in this case, to get traffic data and show it to the driver through the car navigation system.

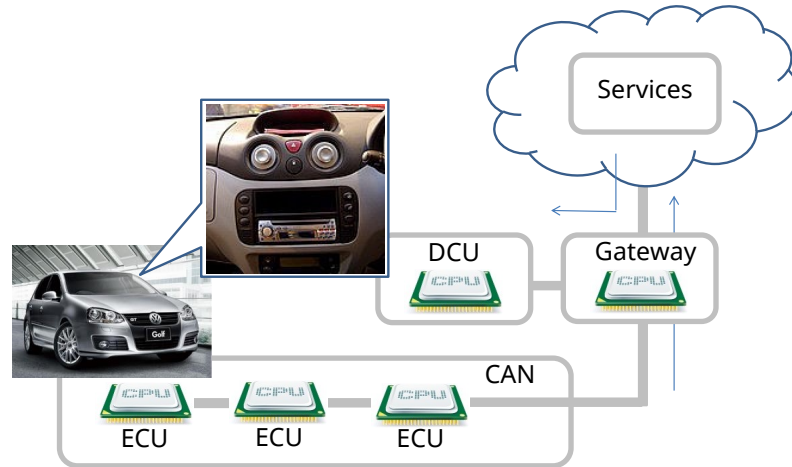


Figure 3 Connected Car

4.2 Common Patterns §

This section introduces common use case patterns that illustrate how devices/things interact with controllers, other devices, agents and servers. In this section, we use the term *client role* as an initiator of a transport protocol, and the term *server role* as a passive component of a transport protocol. This

does not imply prescribing a specific role on any system component. A device can be in a *client* and *server* role simultaneously.

One example of this dual role is a sensor, that registers itself with a cloud service and regularly sends sensor readings to the cloud. In the response messages the cloud can adjust the transmission rate of the sensor's messages or select specific sensor attributes, that are to be transmitted in future messages. Since the sensor registers itself with the cloud and initiates connections, it is in the 'client' role. However, since it also reacts to requests, that are transmitted in response messages, it also fulfills a 'server' role.

The following sections illustrate the roles, tasks, and use case patterns with increasing complexity. They are not exhaustive and are presented to motivate for the WoT architecture and building blocks that are defined in later sections of this specification.

4.2.1 Device Controllers §

The first use case is a local device controlled by a user-operated remote controller as depicted in [Figure 4](#). A remote controller can access an electronic appliance through the local home network directly. In this case, the remote controller can be implemented by a browser or native application.

In this pattern, at least one device like the electronic appliance has a server role that can accept a request from the other devices and responds to them, and sometimes initiates a mechanical action. The other device like the remote controller has a client role that can send a message with a request, like to read a sensor value or to turn on the device. Moreover, to emit a current state or event notification of a device, the device may have a client role that can send a message to another device, which has server roles.




Figure 4 Device Control

4.2.2 Thing-to-Thing §

[Figure 5](#) shows an example of a direct Thing-to-Thing interaction. The scenario is as follows: a sensor detects a change of the room condition, for example the temperature exceeding a threshold, and issues a control message like "turn on" to the electronic appliance. The sensor unit can issue some trigger messages to other devices.

In this case, when two devices that have server roles are connected, at least one device must have also a client role that issues a message to the other to actuate or notify.




Figure 5 Control Agent

4.2.3 Remote Access §

This use case contains a mobile remote controller (e.g., on a smartphone) as shown in [Figure 6](#). The remote controller can switch between different network connections and protocols, e.g., between a cellular network and a home network, which is using protocols such as Wi-Fi and Bluetooth. When the controller is in the home network it is a trusted device and no additional security or access control is required. When it is outside of the trusted network, additional access control and security mechanisms must be applied to ensure a trusted relationship. Note that in this scenario the network connectivity may change due to switching between different network access points or cellular base stations.

In this pattern, the remote controller and the electronic appliance have a client and a server role as in the related scenario in [Figure 4](#).

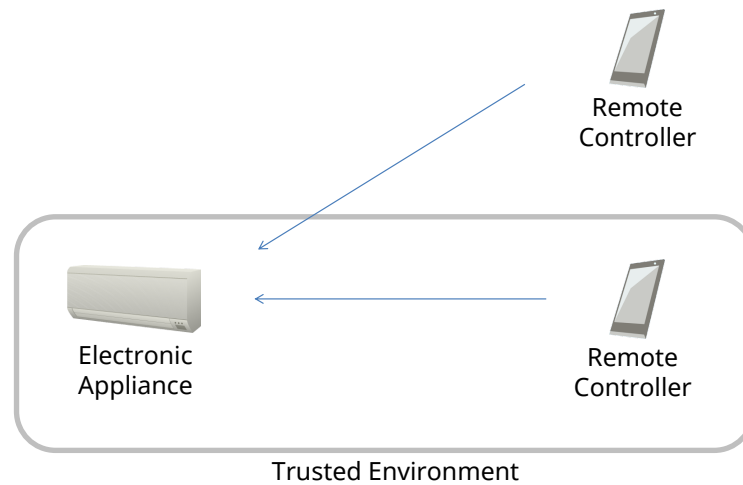


Figure 6 Multiple Network Interfaces

4.2.4 Smart Home Gateways §

[Figure 7](#) shows a use case using a Smart Home Gateway. The gateway is placed between a home network and the Internet. It manages electronic appliances inside the house and can receive commands from a remote controller over the Internet, e.g., from a smartphone as in the previous use case. It is also a virtual representation of a device. The Smart Home Gateway typically offers proxy and firewall functionality.

In this pattern, the home gateway has both a client and a server role. When the remote controller actuates the electronic appliance, it can connect to the electronic appliance in the client role and to the remote controller with the server role. When the electronic appliance emits a message to the remote controller, the gateway act as server roles for the electric appliance, and it act as client roles for the remote controller.




Figure 7 Smart Home Gateway

4.2.5 Edge Devices §

An Edge Device or Edge Gateway is similar to a Smart Home Gateway. We use the term to indicate additional tasks that are carried out by the edge gateway. Whereas the home gateway in [Figure 8](#) primarily just bridges between the public and the trusted network, the edge device has local compute capabilities and typically bridges between different protocols. Edge devices are typically used in industrial solutions, where they can provide preprocessing, filtering and aggregation of data provided by connected devices and sensors.




Figure 8 Edge device

4.2.6 Digital Twins §

A digital twin is a virtual representation, i.e. a model of a device or a group of devices that resides on a cloud server or edge device. It can be used to represent real-world devices which may not be continuously online, or to run simulations of new applications and services, before they get deployed to the real devices.




Figure 9 Digital Twin

Digital twins can model a single device, or they can aggregate multiple devices in a virtual representation of the combined devices.




Figure 10 Digital Twin for Multiple Devices

Digital twins can be realized in different ways, depending on whether a device is already connected to the cloud, or whether it is connected to a gateway, which itself is connected to the cloud.

4.2.6.1 Cloud-ready Devices §

[Figure 11](#) shows an example where electronic appliances are connected directly to the cloud. The cloud mirrors the appliances and, acting as a digital twin, can receive commands from remote controllers (e.g., a smartphone). Authorized controllers can be located anywhere, as the digital twin is globally reachable.




Figure 11 Appliance twin for a Cloud-ready Devices

4.2.6.2 Legacy Devices §

[Figure 12](#) shows an example where legacy electronic appliances cannot directly connect to the cloud. Here, a gateway is needed to relay the connection. The gateway works as:

- integrator of a variety of legacy communication protocols both in the physical and logical view
- firewall toward the Internet
- privacy filter which substitutes real image and/or speech, and logs data locally
- local agent in case the network connection is interrupted
- emergency services running locally when fire alarms and similar events occur

The cloud mirrors the gateway with all connected appliances and acts as a digital twin that manages them in the cloud in conjunction with the gateway. Furthermore, the cloud can receive commands from remote controllers (e.g., a smartphone), which can be located anywhere.




Figure 12 A Digital Twin for a Legacy Device

4.2.7 Multi-Cloud §

Typical IoT deployments consist of multiple (thousands) of devices. Without a standardized mechanism, the management of firmware updates for specific clouds require a lot of effort and hinders wider scale IoT adoption.

The primary benefit of a standardized mechanism for describing devices and device types is the capability of deploying devices to different cloud environments without the need of doing customization at device software / firmware level, i.e., installing cloud specific code to a device. This implies that the solution is flexible enough to describe devices in a way that allows on-boarding and using devices in multiple IoT cloud environments.

This drives adoption of Web of Things devices, since it enables easy usage of new devices in an existing deployment, as well as migration of existing devices from one cloud to the other.

4.2.8 Cross-domain Collaboration §

[Figure 13](#) show an example of a cross-domain collaboration. In this case, each system involves other systems in other domains, such as Smart Factory with Smart City, Smart City with Smart Home. This type of system is called "Symbiotic" ecosystem, as shown in [\[IEC-FOTF\]](#). There are two collaboration

models: direct collaboration and indirect collaboration. In the direct collaboration model, systems exchange information directly with each other in a peer-to-peer manner. In the indirect collaboration, systems exchange information via some collaboration platform. In order to maintain and continue this collaboration, each system provides the metadata of their capabilities and interfaces and adapts itself to others.




Figure 13 Cross-domain collaboration

4.3 Summary §

The previous section described various architecture patterns. In these patterns, some functional entities such as the devices including the legacy devices, controllers, gateways and cloud servers are located at physical locations such as inside building, outside buildings, and data centers. [Figure 14](#) is an overview that shows the combinations and communication paths of these entities.

In a transport protocol layer, each entity arbitrarily selects a suitable role for communications. For example, a device may act as a server when the device provides a service to indefinite number of applications. On the other hand, if a device has limited or intermittent network connectivity, they may act as a client and actively send message to an application when network is available. Regardless of this, in application layer, an application sees that a device provides abstract interfaces to interact and the application can interact with the device using their abstract interfaces.




Figure 14 Use Case Overview

5. Requirements §

This section is normative.

5.1 Functional Requirements §

This section defines the properties required in an abstract Web of Things (WoT) architecture.

5.1.1 Common Principles §

- WoT architecture should enable mutual interworking of different eco-systems using web technology.

- WoT architecture should be based on the web architecture using RESTful APIs.
- WoT architecture should allow to use multiple payload formats which are commonly used in the web.
- WoT architecture must enable different device architectures and must not force a client or server implementation of system components.
- Flexibility
There are a wide variety of physical device configurations for WoT implementations. The WoT abstract architecture should be able to be mapped to and cover all of the variations.
- Compatibility
There are already many existing IoT solutions and ongoing IoT standardization activities in many business fields. The WoT should provide a bridge between these existing and developing IoT solutions and Web technology based on WoT concepts. The WoT should be upwards compatible with existing IoT solutions and current standards.
- Scalability
WoT must be able to scale for IoT solutions that incorporate thousands to millions of devices. These devices may offer the same capabilities even though they are created by different manufacturers.
- Interoperability
WoT must provide interoperability across device and cloud manufacturers. It must be possible to take a WoT enabled device and connect it with a cloud service from different manufacturers out of the box.

5.1.2 Thing Functionalities §

- WoT architecture should allow things to have functionalities such as
 - reading thing's status information
 - updating thing's status information which might cause actuation
 - subscribing to, receiving and unsubscribing to notifications of changes of the thing's status information.
 - invoking functions with input and output parameters which would cause certain actuation or calculation.
 - subscribing to, receiving and unsubscribing to event notifications that are more general than just reports of state transitions.

5.1.3 Search and Discovery §

- WoT architecture should allow clients to know thing's attributes, functionalities and their access points, prior to access to the thing itself.
- WoT architecture should allow clients to search things by its attributes and functionalities.
- WoT architecture should allow semantic search of things providing required functionalities based on a unified vocabulary, regardless of naming of the functionalities.

5.1.4 Description Mechanism §

- WoT architecture should support a common description mechanism which enables describing things and their functions.
- Such descriptions should be not only human-readable, but also machine-readable.
- Such descriptions should allow semantic annotation of its structure and described contents.
- Such description should be able to be exchanged using multiple formats which are commonly used in the web.

5.1.5 Description of Attributes §

- WoT architecture should allow describing thing's attributes such as
 - name
 - explanation
 - version of spec, format and description itself
 - links to other related things and metadata information
- Such descriptions should support internationalization.

5.1.6 Description of Functionalities §

- WoT architecture should allow describing thing's functionalities which is shown in [§ 5.1.2 Thing Functionalities](#)

5.1.7 Network §

- WoT architecture should support multiple web protocols which are commonly used.
- Such protocols include
 1. protocols commonly used in the internet and
 2. protocols commonly used in the local area network
- WoT architecture should allow using multiple web protocols to access to the same functionality.
- WoT architecture should allow using a combination of multiple protocols to the functionalities of the same thing (e.g., HTTP and WebSocket).

5.1.8 Deployment §

- WoT architecture should support a wide variety of thing capabilities such as edge devices with resource restrictions and virtual things on the cloud, based on the same model.
- WoT architecture should support multiple levels of thing hierarchy with intermediate entities such as gateways and proxies.
- WoT architecture should support accessing things in the local network from the outside of the local network (the internet or another local network), considering network address translation.

5.1.9 Application §

- WoT architecture should allow describing applications for a wide variety of things such as edge device, gateway, cloud and UI/UX device, using web standard technology based on the same model.

5.1.10 Legacy Adoption §

- WoT architecture should allow mapping of legacy IP and non-IP protocols to web protocols, supporting various topologies, where such legacy protocols are terminated and translated.
- WoT architecture should allow transparent use of existing IP protocols without translation, which follow RESTful architecture.
- WoT architecture must not enforce client or server roles on devices and services. An IoT device can be either a client or a server, or both, depending on the system architecture; the same is true of edge and cloud services.

5.2 Technical Requirements §

[§ 4.2 Common Patterns](#) defines the Web of Things abstract architecture by showing various use cases and enumerating patterns for combining architectural components. This section describes technical requirements derived from the abstract architecture.

5.2.1 Components in the Web of Things and the Web of Things Architecture §

The use cases help to identify basic components such as devices and applications, that access and control those devices, proxies (i.e., gateways and edge devices) that are located between devices. An additional component useful in some use cases is the directory, which assists with discovery.

Those components are connected to the internet or field networks in offices, factories or other facilities. Note that all components involved may be connected to a single network in some cases, however, in general components can be deployed across multiple networks.

5.2.2 Devices §

Access to devices is made using a description of their functions and interfaces. This description is called *Thing Description (TD)*. A *Thing Description* includes a general metadata about the device, information models representing functions, transport protocol description for operating on information models, and security information.

General metadata contains device identifiers (URI), device information such as serial number, production date, location and other human readable information.

Information models defines device attributes, and represent device's internal settings, control functionality and notification functionality. Devices that have the same functionality have the same information model regardless of the transport protocols used.

Because many systems based on Web of Things architecture are crossing system Domains, vocabularies and meta data (e.g., ontologies) used in information models should be commonly understood by involved parties. In addition to REST transports, PubSub transports are also supported.

Security information includes descriptions about authentication, authorization and secure communications. Devices are required to put TDs either inside them or at locations external to the devices, and to make TDs accessible so that other components can find and access them.

5.2.3 Applications §

Applications need to be able to generate and use network and program interfaces based on metadata (descriptions).

Applications have to be able to obtain these descriptions through the network, therefore, need to be able to conduct search operations and acquire the necessary descriptions over the network.

5.2.4 Digital Twins §

Digital Twins need to generate program interfaces internally based on metadata (descriptions), and to represent virtual devices by using those program interfaces. A twin has to produce a description for the virtual device and make it externally available.

Identifiers of virtual devices need to be newly assigned, therefore, are different from the original devices. This makes sure that virtual devices and the original devices are clearly recognized as separate entities. Transport and security mechanisms and settings of the virtual devices can be different from original devices if necessary. Virtual devices are required to have descriptions provided either directly by the twin or to have them available at external locations. In either case it is required to make the descriptions available so that other components can find and use the devices associated with them.

5.2.5 Discovery §

For TDs of devices and virtual devices to be accessible from devices, applications and twins, there needs to be a common way to share TDs. Directories can serve this requirement by providing functionalities to allow devices and twins themselves automatically or the users to manually register the descriptions.

Descriptions of the devices and virtual devices need to be searchable by external entities. Directories have to be able to process search operations with search keys such as keywords from the general description in the device description or information models.

5.2.6 Security §

Security information related to devices and virtual devices needs to be described in device descriptions. This includes information for authentication/authorization and payload encryption.

WoT architecture should support multiple security mechanism commonly used in the web, such as Basic, Digest, Bearer and OAuth2.0.

5.2.7 Accessibility §

The Web of Things primarily targets machine-to-machine communication. The humans involved are usually developers that integrate Things into applications. End-users will be faced with the front-ends of the applications or the physical user interfaces provided by devices themselves. Both are out of scope of the [W3C](#) WoT specifications. Given the focus on IoT instead of users, accessibility is not a direct requirement, and hence is not addressed within this specification.

There is, however, an interesting aspect on accessibility: Fulfilling the requirements above enables machines to understand the network-facing API of devices. This can be utilized by accessibility tools to provide user interfaces of different modality, thereby removing barriers to using physical devices and IoT-related applications.

6. WoT Architecture §

This section is normative.

To address the use cases in Section 4 and fulfill the requirements in Section 5, the Web of Things (WoT) builds on top of the concept of Web Things – usually simply called Things – that can be used by so-called Consumers. This section provides the background and normative assertions to define the overall [W3C](#) Web of Things architecture. As the Web of Things addresses stakeholders from different domains, certain aspects of Web technology are explained in more detail, in particular the concept of hypermedia.

6.1 Overview §

A Thing is the abstraction of a physical or virtual entity (e.g., a device or a room) and is described by standardized metadata. In [W3C](#) WoT, the description metadata *MUST* be a WoT Thing Description (TD) [[WOT-THING-DESCRIPTION](#)]. Consumers *MUST* be able to parse and process the TD representation format, which is based on JSON [[RFC8259](#)]. The format can be processed either through classic JSON libraries or a JSON-LD processor, as the underlying information model is graph-based and its serialization compatible with JSON-LD 1.1 [[JSON-LD11](#)]. The use of a JSON-LD processor for processing a TD additionally enables semantic processing including transformation to RDF triples, semantic inference and accomplishing tasks given based on ontological terms, which would make Consumers behave more autonomous. A TD is instance-specific (i.e., describes an individual Thing, not types of Things) and is the default external, textual (Web) representation of a Thing. There *MAY* be other representations of a Thing such as an HTML-based user interface, simply an image of the physical entity, or even non-Web representations in closed systems.

To be a Thing, however, at least one TD representation *MUST* be available. The WoT Thing Description is a standardized, machine-understandable representation format that allows Consumers to discover and interpret the capabilities of a Thing (through semantic annotations) and to adapt to different implementations (e.g., different protocols or data structures) when interacting with a Thing, thereby enabling interoperability across different IoT platforms, i.e., different ecosystems and standards.



Figure 15 Consumer-Thing interaction

A Thing can also be the abstraction of a virtual entity. A virtual entity is the composition of one or more Things (e.g., a room consisting of several sensors and actuators). One option for the composition is to provide a single, consolidated WoT Thing Description that contains the superset of capabilities for the virtual entity. In cases where the composition is rather complex, its TD may *link* to hierarchical sub-Things within the composition. The main TD acts as entry point and only contain general metadata and potentially overarching capabilities. This allows grouping of certain aspects of more complex Things.

Linking does not only apply to hierarchical Things, but relations between Things and other resources in general. Link relation types express how Things relate, for instance, a switch controlling a light or a room monitored by a motion sensor. Other resources related to a Thing can be manuals, catalogs for spare parts, CAD files, a graphical UI, or any other document on the Web. Overall, Web linking among Things makes the Web of Things navigable, for both humans and machines. This can be further facilitated by providing Thing directories that manage a catalog of available Things, usually by caching their TD representation. In summary, WoT Thing Descriptions *MAY* link to other Things and other resources on the Web to form a Web of Things.




Figure 16 Linked Things

Things must be hosted on networked system components with a software stack to realize interaction through a network-facing interface, the WoT Interface of a Thing. One example of this is an HTTP server running on an embedded device with sensors and actuators interfacing the physical entity behind the Thing abstraction. However, W3C WoT does not mandate where Things are hosted; it can be on the IoT device directly, an Edge device such as a gateway, or the cloud.

A typical deployment challenge is a scenario, where local networks are not reachable from the Internet, usually because of IPv4 Network Address Translation (NAT) or firewall devices. To remedy this situation, W3C WoT allows for Intermediaries between Things and Consumers.

Intermediaries can act as proxies for Things, where the Intermediary has a WoT Thing Description similar to the original Thing, but which points to the WoT Interface provided by the Intermediary. Intermediaries may also augment existing Things with additional capabilities or compose a new Thing out of multiple available Things, thereby forming a virtual entity. To Consumers, Intermediaries look like Things, as they possess WoT Thing Descriptions and provide a WoT Interface, and hence might be indistinguishable from Things in a layered system architecture like the Web [REST]. An identifier in the WoT Thing Description *MUST* allow for the correlation of multiple TDs representing the same original Thing or ultimately unique physical entity.




Figure 17 Intermediary

Another remedy for restricted local networks is binding the WoT Interface to a protocol that establishes the connection from the Thing within the local network to a publicly reachable Consumer.

Things *MAY* be bundled together with a Consumer to enable Thing-to-Thing interaction. Usually, the Consumer behavior is embedded in the software component, which is also implementing the behavior of the Thing. The configuration of the Consumer behavior *MAY* be exposed through the Thing.

The concepts of W3C WoT are applicable to all levels relevant for IoT applications: the device level, edge level, and cloud level. This fosters common interfaces and APIs across the different levels and enables various integration patterns such as Thing-to-Thing, Thing-to-Gateway, Thing-to-Cloud, Gateway-to-Cloud, and even cloud federation, i.e., interconnecting cloud computing environments of two or more service providers, for IoT applications. [Figure 18](#) gives an overview how the WoT concepts introduced above can be applied and combined to address the use cases summarized in [§ 4.3 Summary](#).

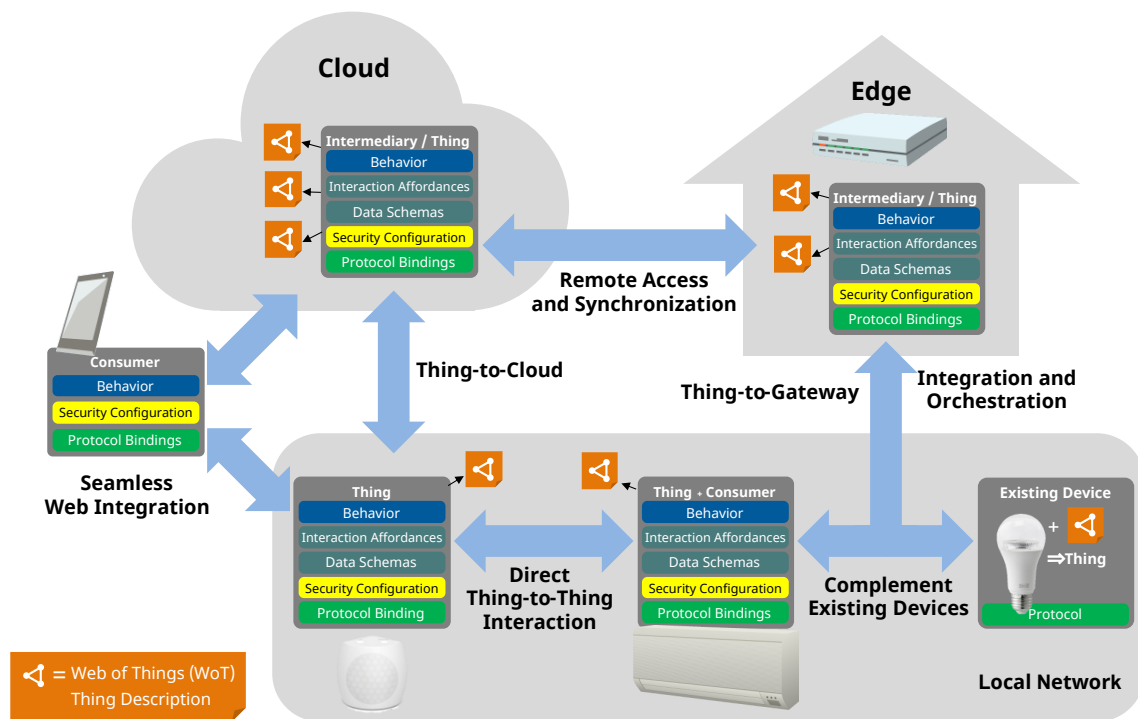


Figure 18 Abstract Architecture of W3C WoT

6.2 Affordances §

A central aspect in W3C WoT is the provision of machine-understandable metadata (i.e., [WoT Thing Descriptions](#)). Ideally, such metadata is self-descriptive, so that [Consumers](#) are able to identify *what* capabilities a [Thing](#) provides and *how* to use the provided capabilities. A key to this self-descriptiveness lies in the concept of affordances.

The term affordance originates in ecological psychology, but was adopted in the field of Human-Computer Interaction [HCI] based on the definition by Donald Norman: "'Affordance' refers to the perceived and actual properties of the thing, primarily those fundamental properties that determine just how the thing could possibly be used." [NORMAN]

An example for this is a door with a handle. The door handle is an affordance, which suggests that the door can be opened. For humans, a door handle usually also suggests *how* the door can be opened; an American knob suggests twisting, a European lever handle suggests pressing down.

The hypermedia principle, which is one of the core foundations of the REST architectural style [REST], demands that any piece of information available on the Web be linked to other pieces of information so that the consumer of the information gets explicit knowledge about how to navigate the Web and control Web applications. Here, the simultaneous presentation of information and control (provided in the form of hyperlinks) is a mechanism that *affords* Web clients the means to drive Web applications. In this context, an affordance is the description of a hyperlink (e.g., via a link relation type and link target attributes) suggesting Web clients how to navigate and possibly how to act on the linked resource. Hence, links provide navigation affordances.

Drawn from this hypermedia principle, the Web of Things defines Interaction Affordances as metadata of a Thing that shows and describes the possible choices to Consumers, thereby suggesting how Consumers may interact with the Thing. A general Interaction Affordance is navigation, which is activated by following a link, thereby enabling Consumers to browse the Web of Things. § 6.4 Interaction Model defines three more types of Interaction Affordances for W3C WoT: Properties, Actions, and Events.

Overall, this W3C WoT definition is aligned with HCI and interaction designers, who create physical Things, as well as the REST and microservice community, who is working on Web services in general.

6.3 Web Thing §

A Web Thing has four architectural aspects of interest: its *behavior*, its Interaction Affordances, its *security configuration*, and its Protocol Bindings, as depicted in Figure 19. The behavior aspect of a Thing includes both the autonomous behavior and the handlers for the Interaction Affordances. The Interaction Affordances provide a model of how Consumers can interact with the Thing through abstract operations, but without reference to a specific network protocol or data encoding. The protocol binding adds the additional detail needed to map each Interaction Affordance to concrete messages of a certain protocol. In general, different concrete protocols may be used to support different subsets of Interaction Affordances, even within a single Thing. The security configuration

aspect of a Thing represents the mechanisms used to control access to the Interaction Affordances and the management of related Public Security Metadata and Private Security Data.

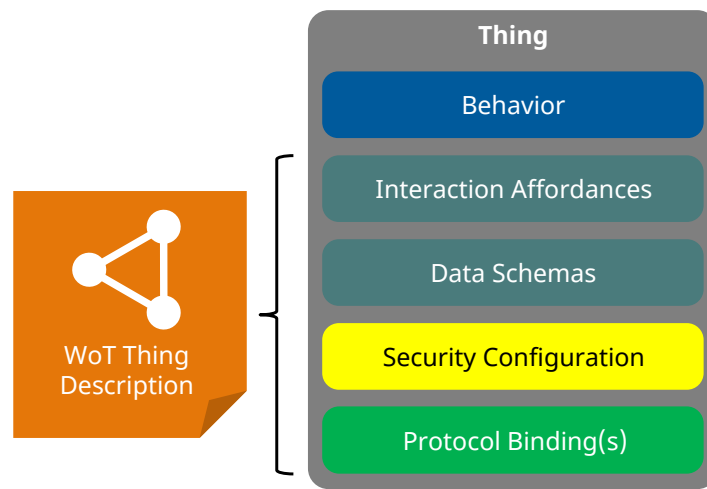


Figure 19 Architectural Aspects of a Thing

6.4 Interaction Model §

Originally, a Web resource usually represented a document on the World Wide Web that can simply be fetched by a Web client. With the introduction of Web services, resources became more generic interaction entities that can implement any kind of behavior. This very high level of abstraction makes it hard to provide a loose coupling between applications and resources due to the manifold interaction possibilities. As a result, at the time of writing typical API descriptions consist of a static mapping from an application intent to a resource address, method, request payload structure, response payload structure, and expected errors. This imposes a tight coupling between Web client and Web service.

The Interaction Model of W3C WoT introduces an intermediate abstraction that formalizes the mapping from application intent to concrete protocol operations and also narrows the possibilities how Interaction Affordances can be modeled.

In addition to navigation affordances (i.e., Web links), Things *MAY* offer three other types of Interaction Affordances defined by this specification: Properties, Actions, and Events. While this narrow waist allows to decouple Consumers and Things, these four types of Interaction Affordances are still able to model virtually all interaction possibilities found in IoT devices and services.

6.4.1 Properties §

A Property is an Interaction Affordance that exposes the state of the Thing. The state exposed by a Property *MUST* be retrievable (readable). Optionally, the state exposed by a Property *MAY* be updated (writeable). Things *MAY* choose to make Properties observable by pushing the new state after a change (cf. Observing Resources [[RFC7641](#)]). Write-only state should be updated through an Action.

If the data is not fully specified by the Protocol Binding used (e.g., through a media type), Properties *MAY* contain one data schema for the exposed state.

Examples of Properties are sensor values (read-only), stateful actuators (read-write), configuration parameters (read-write), Thing status (read-only or read-write), or computation results (read-only).

6.4.2 Actions §

An Action is an Interaction Affordance that allows to invoke a function of the Thing. An Action *MAY* manipulate state that is not directly exposed (cf. Properties), manipulate multiple Properties at a time, or manipulate Properties based on internal logic (e.g., toggle). Invoking an Action *MAY* also trigger a process on the Thing that manipulates state (including physical state through actuators) over time.

If the data is not fully specified by the Protocol Binding used (e.g., through a media type), Actions *MAY* contain data schemas for optional input parameters and output results.

Examples of Actions are changing multiple Properties simultaneously, changing Properties over time such as fading the brightness of a light (dimming) or with a process that shall not be disclosed such as a proprietary control loop algorithm, or invoking a long-lasting process such as printing a document.

6.4.3 Events §

An Event Interaction Affordance describes an event source that pushes data asynchronously from the Thing to the Consumer. Here not state, but state transitions (i.e., events) are communicated. Events *MAY* be triggered through conditions that are not exposed as Properties.

If the data is not fully specified by the Protocol Binding used (e.g., through a media type), Events *MAY* contain data schemas for the event data and possible subscription control messages (e.g., to subscribe with a Webhook callback URI).

Examples of Events are discrete events such as an alarm or samples of a time series that are pushed regularly.

6.5 Hypermedia Controls §

On the Web, an affordance is the simultaneous presentation of information and controls, such that the information becomes the affordance through which the user obtains choices. For humans, the information is usually text or images describing or decorating a hyperlink. The control is a Web link, which includes at least the URI of the target resource, which can be dereferenced by the Web browser (i.e., the link can be followed). But also machines can follow links in a meaningful way, when the Web link is further described by a relation type and a set of target attributes. A hypermedia control is the machine-understandable description of *how* to activate an affordance. Hypermedia controls usually originate from a Web server and are discovered in-band while a Web client is interacting with the server. This way, Web servers can drive clients through Web applications dynamically, by taking their current state and other factors such as authorization into account. This is opposed to out-of-band interface descriptions that need to be preinstalled or hardcoded into clients (e.g., RPC, WS-* Web services, HTTP services with fixed URI-method-response definitions).

W3C WoT makes use of two kinds of hypermedia controls: *Web links* [RFC8288], the well-established control to navigate the Web, and Web forms as a more powerful control to enable any kind of operation. Links are already used in other IoT standards and IoT platforms such as CoRE Link Format [RFC6690], OMA LWM2M [LWM2M], and OCF [OCF]. Form is a new concept that besides W3C WoT is also introduced by the *Constrained RESTful Application Language (CoRAL)* [CoRAL] defined by the IETF.

6.5.1 Links §

Links enable Consumers (or Web clients in the broader sense) to change the current context (cf. the set of resource representations currently rendered in the Web browser) or to include additional resources into the current context, depending on the relation between context and link target. Consumers do so by *dereferencing* the target URI, that is, fetching the resource representation by following a link.

W3C WoT follows the definitions of Web Linking [RFC8288], where a link is comprised of:

- a link context,
- a relation type,
- a link target, and
- optionally target attributes.

Link relation types are either a set of predefined tokens that are registered with IANA [IANA-RELATIONS], which must adhere to the ABNF [RFC5234]

`LOALPHA *(LOALPHA / DIGIT / "." / "-")` (e.g., `stylesheet`), or extension types in the

form of URIs [\[RFC3986\]](#). Extension relation types *MUST* be compared as strings using a case-insensitive comparison. (If they are serialized in a different format they are to be converted to URIs). Nevertheless, all-lowercase URIs *SHOULD* be used for extension relation types. [\[RFC8288\]](#)

In the Web of Things, links are used for discovery and to express relations between Things (e.g., hierarchical or functional) and relations to other documents on the Web (e.g., manuals or alternative representations such as CAD models).

6.5.2 Forms §

Forms enable Consumers (or Web clients in the broader sense) to perform operations that go beyond dereferencing a URI (e.g., to manipulate the state of a Thing). Consumers do so by *filling out* and *submitting* the form to its submission target. This usually requires more detailed information about the contents of the (request) message than a link can provide (e.g., method, header fields, or other protocol options). Forms can be seen as a request template, where the provider pre-filled parts of the information according to its own interface and state, and left parts blank to be filled by the Consumers (or Web client in general).

W3C WoT defines forms as new hypermedia control. Note that the definition in CoRAL is virtually identical, and hence compatible [\[CoRAL\]](#). A form is comprised of:

- a form context,
- an operation type,
- a submission target,
- a request method, and
- optionally form fields.

A form can be viewed as a statement of "To perform an **operation type** operation on **form context** , issue a **request method** request to **submission target** " where the optional form fields may further describe the required request.

Form contexts and submission targets *MUST* both be Internationalized Resource Identifiers (IRIs) [\[RFC3987\]](#). However, in the common case, they will also be URIs [\[RFC3986\]](#), because many protocols (such as HTTP) do not support IRIs.

Form context and submission target *MAY* point to the same resource or different resources, where the submission target resource implements the operation for the context.

The operation type identifies the semantics of the operation. Operation types are denoted similar to link relation types:

- Well-known operation types *MUST* follow the ABNF `LOALPHA *(LOALPHA / DIGIT / "." / "-")`. Well-known operation types *MUST* be compared using a case-insensitive comparison. The well-known operation types for the Web of Things defined by this specification are given in [Table 1](#).
- The set of predefined operation types *MAY* be augmented by *Extension operation types* chosen by applications. Extension operation types *MUST* be URIs [\[RFC3986\]](#) that uniquely identify the type. Extension operation types *MUST* be compared as strings using a case-insensitive comparison. Nevertheless, all-lowercase URIs *SHOULD* be used for extension operation types.

The request method *MUST* identify one method of the standard set of the protocol identified by the submission target URI scheme.

Form fields are optional and *MAY* further specify the expected request message for the given operation. Note that this is not limited to the payload, but may affect also protocol headers. Form fields *MAY* depend on the protocol used for the submission target as specified in the URI scheme. Examples are HTTP header fields, CoAP options, the protocol-independent media type [\[RFC2046\]](#) including parameters (i.e., full content type) for the request payload, or information about the expected response.

Table 1 Well-known Operation Types for the Web of Things

<i>Operation Type</i>	<i>Description</i>
readproperty	Identifies the read operation on Property Affordances to retrieve the corresponding data.
writeproperty	Identifies the write operation on Property Affordances to update the corresponding data.
observeproperty	Identifies the observe operation on Property Affordances to be notified with the new data when the Property was updated.
unobserveproperty	Identifies the unobserve operation on Property Affordances to stop the corresponding notifications.
invokeaction	Identifies the invoke operation on Action Affordances to perform the corresponding action.

<i>Operation Type</i>	<i>Description</i>
subscribeevent	Identifies the subscribe operation on Event Affordances to be notified by the Thing when the event occurs.
unsubscribeevent	Identifies the unsubscribe operation on Event Affordances to stop the corresponding notifications.
readallproperties	Identifies the readallproperties operation on Things to retrieve the data of all Properties in a single interaction.
writeallproperties	Identifies the writeallproperties operation on Things to update the data of all writable Properties in a single interaction.
readmultipleproperties	Identifies the readmultipleproperties operation on Things to retrieve the data of selected Properties in a single interaction.
writemultipleproperties	Identifies the writemultipleproperties operation on Things to update the data of selected writable Properties in a single interaction.

EDITOR'S NOTE

As of this specification, the well-known operation types are a fixed set that results from the [WoT Interaction Model](#). Other specifications may define further well-known operation types that are valid for their respective document format or form serialization. Later versions of this specification or another specification may set up an IANA registry in the future to enable extension and a more generic Web form model that may be applied beyond WoT specifications.

6.6 Protocol Bindings §

A Protocol Binding is the mapping from an [Interaction Affordance](#) to concrete messages of a specific protocol such as HTTP [[RFC7231](#)], CoAP [[RFC7252](#)], or MQTT [[MQTT](#)]. It informs the [Consumer](#) *how* to activate the [Interaction Affordance](#) through a network-facing interface. The [Protocol Bindings](#) follow the Uniform Interface constraint of REST [[REST](#)] to support interoperability. Thus, not all

communication protocols are eligible to implement [Protocol Bindings](#) for [W3C WoT](#); the requirements are given in the assertions below.

In the door example given in [§ 6.2 Affordances](#), the [Protocol Binding](#) corresponds to the door handle at the level of knob vs lever, which suggests *how* the door can be opened.

6.6.1 Hypermedia-driven §

Interaction Affordances *MUST* include one or more Protocol Bindings. Protocol Bindings *MUST* be serialized as hypermedia controls (see [§ 6.5 Hypermedia Controls](#)) to be self-descriptive on how to activate the Interaction Affordance. The hypermedia controls *MUST* originate from the authority managing the Thing that is providing the corresponding Interaction Affordance. The authority can be the [Thing](#) itself, producing the [TD](#) document at runtime (based on its current state and including network parameters such as its IP address) or serving it from memory with only the current network parameters inserted. The authority can also be an external entity that has full and up-to-date knowledge of the [Thing](#) including its network parameters and internal structure (e.g., software stack). This enables a loose coupling between [Things](#) and [Consumers](#), allowing for an independent lifecycle and evolution. The hypermedia controls *MAY* be cached outside the [Thing](#) and used for offline processing if caching metadata is available to determine the freshness.

6.6.2 URIs §

Eligible protocols for [W3C WoT](#) *MUST* have an associated URI scheme [\[RFC3986\]](#) that is registered with IANA (see [\[IANA-URI-SCHEMES\]](#)). Hypermedia controls rely on URIs [\[RFC3986\]](#) to identify link and submission targets. Thereby, the URI scheme (the first component up to ":",") identifies the communication protocol to be used for [Interaction Affordances](#) with the Thing. [W3C WoT](#) refers to these protocols as [transfer protocols](#).

6.6.3 Standard Set of Methods §

Eligible protocols for [W3C WoT](#) *MUST* be based on a standard set of methods that are known a priori. The standard set of methods makes messages self-descriptive to enable intermediate processing of [Interaction Affordances](#), for instance by proxies or to translate between Protocol Bindings [\[REST\]](#). Furthermore, it allows [Consumers](#) to have re-usable protocol stacks of common [transfer protocols](#) such as HTTP, CoAP, or MQTT, avoiding Thing-specific code or plugins for [Consumers](#).

6.6.4 Media Types §

All data (a.k.a. content) exchanged when activating Interaction Affordances *MUST* be identified by a media type [RFC2046] in the Protocol Binding. Media types are labels to identify representation formats, for instance `application/json` for JSON [RFC8259] or `application/cbor` for CBOR [RFC7049]. They are managed by IANA.

Some media types might need additional parameters to fully specify the representation format used. Examples are `text/plain; charset=utf-8` or `application/ld+json; profile="http://www.w3.org/ns/json-ld#compacted"`. This needs to be considered in particular when describing data to be sent to Things. There might also be standardized transformations on the data such as content coding [RFC7231]. Protocol Bindings *MAY* have additional information that specifies representation formats in more detail than the media type alone.

Note that many media types only identify a generic serialization format that does not provide further semantics for its elements (e.g., XML, JSON, CBOR). Thus, the corresponding Interaction Affordances *SHOULD* declare a data schema to provide more detailed syntactic metadata for the data exchanged.

6.7 WoT System Components and Their Interconnectivity §

Section § 6.1 Overview described the WoT architecture in terms of the abstract WoT architecture components such as Things, Consumers and Intermediaries. When those abstract WoT architecture components are implemented as a software stack to take a specific role in the WoT architecture, such software stacks are called Servients. Systems that are based on the WoT architecture involve Servients, which are communicating with each other to achieve the goals of a system.

This section uses system configuration diagrams to illustrate how Servients work together to build systems based on the WoT architecture.

A Thing can be implemented by a Servient. In a Thing, a Servient software stack contains a representation of a Thing called Exposed Thing, and makes its WoT Interface available to Consumers of the Thing. This Exposed Thing may be used by other software components on the Servient (e.g., applications) to implement the behavior of the thing.




Figure 20 Servient as a Thing

On the other hand, Consumers are always implemented by Servients, as they must be able to process the Thing Description (TD) format and must have a protocol stack that can be configured through Protocol Binding information contained in the TDs.

In a Consumer, a Servient software stack provides a representation of a Thing called Consumed Thing, and makes it available to those applications running on the Servient that need to process TDs to interact with Things.




Figure 21 Servient as a Consumer

A Consumed Thing instance in the Servient software stack serves to separate the protocol level complexity from applications. It is communicating with Exposed Things on behalf of the application.

Similarly, an Intermediary is yet another WoT architecture component implemented by a Servient. An Intermediary is located between a Thing and its Consumers, performing the roles of both a Consumer (to the Thing) and a Thing (to the Consumers). In an Intermediary, a Servient software stack contains the representations of both a Consumer (Consumed Thing) and a Thing (Exposed Thing).




Figure 22 Servient as an Intermediary

6.7.1 Direct Communication §

[Figure 23](#) shows direct communication between a [Thing](#), which is exposing [Interaction Affordances](#) through [Thing Descriptions](#), and a [Consumer](#) that uses the [Thing](#) by means of the [Interaction Affordances](#). Direct communication applies when both [Servients](#) use the same network protocol(s) and are accessible to each other.



Figure 23 High-level architecture of Consumer and Thing

An [Exposed Thing](#) is the software representation of a [Thing](#) abstraction, serving a [WoT Interface](#) of the [Interaction Affordances](#) provided by the [Thing](#).

A [Consumed Thing](#) is the software representation of a remote [Thing](#) being consumed by a [Consumer](#), serving as the interface to the remote [Thing](#) for the applications. A [Consumer](#) can generate a [Consumed Thing](#) instance by parsing and processing a [TD](#) document. Interactions between a [Consumer](#) and a [Thing](#) are performed by the [Consumed Thing](#) and the [Exposed Thing](#) exchanging messages over a direct network connection between them.

6.7.2 Indirect Communication §

In [Figure 24](#), a [Consumer](#) and a [Thing](#) connect to each other via an [Intermediary](#). An [Intermediary](#) is required if the [Servients](#) use different protocols or if they are on different networks that require authentication and provide access control (e.g. firewalls).




Figure 24 High-level architecture with Intermediary

An [Intermediary](#) combines [Exposed Thing](#) and [Consumed Thing](#) functionality. The functionality of [Intermediaries](#) includes relaying messages for the [Interaction Affordances](#) between a [Consumer](#) and a

Thing, optionally caching the Thing's data for faster response, and transforming communication when the functionality of the Thing is extended by the Intermediary. In an Intermediary, a Consumed Thing creates a proxy object of the Exposed Thing of a Thing, and a Consumer can access the proxy object (i.e., the Exposed Thing of the Intermediary) through its own Consumed Thing.

Consumer and Intermediary can communicate in a different protocol than Intermediary and Thing. For example, an Intermediary can provide a bridge between a Thing that uses CoAP and the application of a Consumer that uses HTTP.

Even when there are multiple different protocols used between Intermediary and Things, Consumer can indirectly communicate with those Things using a single protocol through the Intermediary. The same is true for the authentication. The Consumed Thing of a Consumer only needs to authenticate with the Exposed Things of the Intermediary using a single security mechanism, while the Intermediary might need multiple security mechanism to authenticate with different Things.

Usually, an Intermediary generates the Thing Description for its proxy object based on the Thing Description of the originating Thing. Depending on the requirements of the use cases, the TD for the proxy object may either use the same identifier as the TD of the original Thing, or it gets assigned a new identifier. If necessary, a TD generated by an Intermediary *MAY* contain interfaces for other communication protocols.

7. WoT Building Blocks §

This section is normative.

The Web of Things (WoT) building blocks allow the implementation of systems that conform with the abstract WoT Architecture. The specifics of these building blocks are defined in separate specification; this section provides an overview and a summary.

The WoT building blocks support each of the architectural aspects of a Thing discussed in § 6.3 Web Thing and depicted in Figure 19. The individual building blocks are shown in the context of an abstract Thing in Figure 25. This is an abstract view and does not represent any particular implementation; instead it illustrates the relationship between the building blocks and the main architectural aspects of a Thing. In this figure the WoT building blocks are highlighted with black outlines. Security, a cross-cutting concern, is separated into public and protected private components. The WoT Scripting API is optional and the Binding Templates are informative.

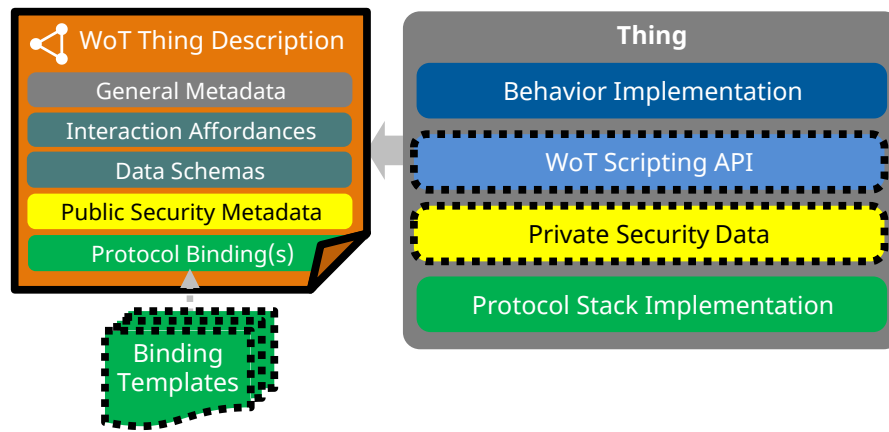


Figure 25 Relationship of WoT Building Blocks to the Architectural Aspects of a Thing.

In the following sections we will provide additional information on each WoT building block: the [WoT Thing Description](#), the [WoT Binding Templates](#), and the [WoT Scripting API](#). Security, although it is a cross-cutting concern, can be considered a fourth building block.

7.1 WoT Thing Description §

The [WoT Thing Description](#) (TD) specification [[WOT-THING-DESCRIPTION](#)] defines an *information model* based on a semantic vocabulary and a *serialized representation based on JSON*. TDs provide rich metadata for [Things](#) in a way that is both human-readable and machine-understandable. Both the information model and the representation format of [TDs](#) are aligned with [Linked Data](#) [[LINKED-DATA](#)], so that besides raw JSON processing, implementations may choose to make use of JSON-LD [[JSON-LD11](#)] and graph databases to enable powerful semantic processing of the metadata.

A [Thing Description](#) (TD) describes [Thing](#) instances with general metadata such as name, ID, descriptions, and also can provide relation metadata through links to related [Things](#) or other documents. [TDs](#) also contain [Interaction Affordance](#) metadata based on the interaction model defined in § 6.4 [Interaction Model](#); [Public Security Metadata](#); and communications metadata defining [Protocol Bindings](#). The [TD](#) can be seen as the *index.html for Things*, as it provides the entry point to learn about the provided services and related resources, both of which are described using hypermedia controls.

Ideally, the [TD](#) is created and/or hosted by the [Thing](#) itself and retrieved upon discovery. Yet it can also be hosted externally when a [Thing](#) has resource restrictions (e.g., limited memory space, limited power) or when an existing device is retrofitted to become part of the Web of Things. A common pattern to improve discovery (e.g., for constrained devices) and to facilitate device management is to register [TDs](#) with a directory. It is recommended that Consumers use a [TD](#) caching mechanism

combined with a notification mechanism, which will inform them when it is required to fetch a new version of the TD, in case the Thing is updated.

For semantic interoperability, TDs may make use of a domain-specific vocabulary, for which explicit extension points are provided. However, development of any particular domain-specific vocabulary is currently out-of-scope of the W3C WoT standardization activity.

Three examples of potentially useful external IoT vocabularies are SAREF [[SAREF](#)], Schema Extensions for IoT [[IOT-SCHEMA-ORG](#)], and the W3C Semantic Sensor Network ontology [[VOCAB-SSN](#)]. Use of such external vocabularies in TDs is optional. In the future additional domain-specific vocabularies may be developed and used with TDs.

Overall, the WoT Thing Description building block fosters interoperability in two ways: First, TDs enable machine-to-machine communication in the Web of Things. Second, TDs can serve as a common, uniform format for developers to document and retrieve all the details necessary to create applications that can access IoT devices and make use of their data.

7.2 WoT Binding Templates §

This section is non-normative.

The IoT uses a variety of protocols for accessing devices, since no single protocol is appropriate in all contexts. Thus, a central challenge for the Web of Things is to enable interactions with the plethora of different IoT platforms (e.g., OCF, oneM2M, OMA LWM2M, OPC UA) and devices that do not follow any particular standard, but provide an eligible interface over a suitable network protocol. WoT is tackling this variety through Protocol Bindings, which must meet a number of constraints (see [§ 6.6 Protocol Bindings](#)).

The non-normative WoT Binding Templates specification [[WOT-BINDING-TEMPLATES](#)] provides a collection of communication metadata blueprints that give guidance on how to interact with different IoT platforms. When describing a particular IoT device or service, the Binding Template for the corresponding IoT Platform can be used to look up the communication metadata that must be provided in the Thing Description to support that platform.

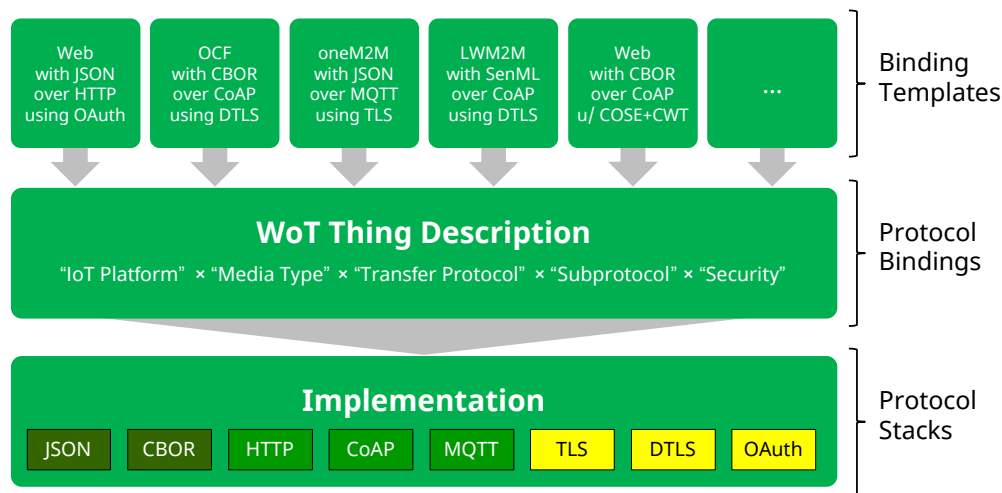


Figure 26 From Binding Templates to Protocol Bindings

Figure 26 shows how Binding Templates are applied. A WoT Binding Template is created only once for each IoT Platform and can then be reused in all TDs for devices of that platform. The Consumer that is processing a TD must implement the required Protocol Binding by including a corresponding protocol stack and by configuring the stack (or its messages) according to the information given in the TD.

The communication metadata of Protocol Bindings spans five dimensions:

- **IoT Platform:**

IoT Platforms often introduce proprietary modifications at the application layer such as platform-specific HTTP header fields or CoAP options. Forms (see § 6.5.2 Forms) may contain the necessary information to apply these tweaks in additional form fields defined for the application-layer protocol used.

- **Media Type:**

IoT Platforms often differ in the representation formats (a.k.a. serializations) used for exchanging data. The media type [RFC2046] identifies these formats, while media type parameters may specify them further. Forms may contain the media type and optional parameters in additional form fields such as a content type field known from HTTP, which combines media type and its potential parameters (e.g., `text/plain; charset=utf-8`).

- **Transfer Protocol:**

The Web of Things uses the term transfer protocol for the underlying, standardized application-layer protocol without application-specific options or subprotocol mechanisms. The URI scheme

of the form submission target contains the information required to identify the transfer protocol, e.g., HTTP, CoAPS, or WebSocket through **http:**, **coaps:**, or **ws:**, resp.

- **Subprotocol:**

Transfer protocols may have extension mechanisms that must be known to interact successfully. Such subprotocols cannot be identified from the URI scheme alone and must be declared explicitly. Examples are the push notification workarounds for HTTP such as long polling [[RFC6202](#)] or Server-Sent Events [[HTML](#)]. Forms may contain the necessary information to identify the subprotocol in additional form fields.

- **Security:**

Security mechanisms can be applied at different layers of the communication stack and might be used together, often to complement each other. Examples are (D)TLS [[RFC8446](#)]/[[RFC6347](#)], IPSec [[RFC4301](#)], OAuth [[RFC6749](#)], and ACE [[RFC7744](#)]. Due to the cross-cutting nature of security, the necessary information to apply the right mechanism may be given within the general metadata of the Thing and/or specialized for each Interaction Affordance or form.

7.3 WoT Scripting API §

This section is non-normative.

The WoT Scripting API is an optional "convenience" building block of W3C WoT that eases IoT application development by providing an ECMAScript-based API [[ECMAScript](#)] similar to the Web browser APIs. By integrating a scripting runtime system into the WoT Runtime, the WoT Scripting API enables using portable application scripts that define the behavior of Things, Consumers, and Intermediaries.

Traditionally, IoT device logic is implemented in firmware, which results in productivity constraints similar to that of embedded development, including a relatively complex update process. The WoT Scripting API in contrast enables implementing device logic by reusable scripts executed in a runtime system for IoT applications not dissimilar to that of a Web browser, and aims to improve productivity and reduce integration costs. Furthermore, standardized APIs enable portability for application modules, for instance, to move compute-intense logic from a device up to a local gateway, or to move time-critical logic from the cloud down to a gateway or edge node.

The non-normative WoT Scripting API specification [[WOT-SCRIPTING-API](#)] defines the structure and algorithms of the programming interface that allows scripts to discover, fetch, consume, produce, and expose WoT Thing Descriptions. The runtime system of the WoT Scripting API instantiates local

objects that act as an interface to other Things and their Interaction Affordances (Properties, Actions, and Events). It also allows scripts to expose Things, that is, to define and implement Interaction Affordances and publish a Thing Description.

7.4 WoT Security and Privacy Guidelines §

This section is non-normative.

Security is a cross-cutting concern and should be considered in all aspects of system design. In the WoT architecture, security is supported by certain explicit features, such as support for Public Security Metadata in TDs and by separation of concerns in the design of the WoT Scripting API. The specification for each building block also includes a discussion of particular security and privacy considerations of that building block. Another non-normative specification, the *WoT Security and Privacy Guidelines* [[WOT-SECURITY](#)], provides additional cross-cutting security and privacy guidance.

8. Abstract Servient Architecture §

This section is non-normative.

As defined in § 6.7 WoT System Components and Their Interconnectivity, a Servient is a software stack that implements the WoT building blocks presented in the previous section. Servients can host and expose Things and/or consume Things (i.e., host Consumers). Depending on the Protocol Binding, Servients can perform in both server and client role, hence the portmanteau naming.

The previous section describes how the WoT building blocks conceptually relate to each other and how they correspond to the abstract WoT Architecture (see § 6. WoT Architecture). When implementing these concepts, a more detailed view is necessary that takes certain technical aspects into account. This section describes the detailed architecture of a Servient implementation.

[Figure 27](#) shows a Servient implementation that is using the (optional) WoT Scripting API building block. Here, the WoT Runtime is also a Scripting Runtime system that, in addition to managing the WoT-specific aspects, also interprets and executes the application scripts. Servients that support the WoT Scripting API usually run on powerful devices, edge nodes, or in the cloud. The WoT Architecture does not limit the application-facing API of the WoT Runtime to JavaScript/ECMAScript. Also other runtime systems can be used to implement a Servient.

Section § 8.8.1 Native WoT API presents an alternative Servient implementation without the WoT Scripting API building block. The WoT Runtime may use any programming language for its

application-facing API. Usually, it is the native language of the Servient software stack, for instance C/C++ for embedded Servients or Java for cloud-based Servients. It may also be an alternative scripting language such as Lua to combine the benefits of application scripts with low resource consumption.

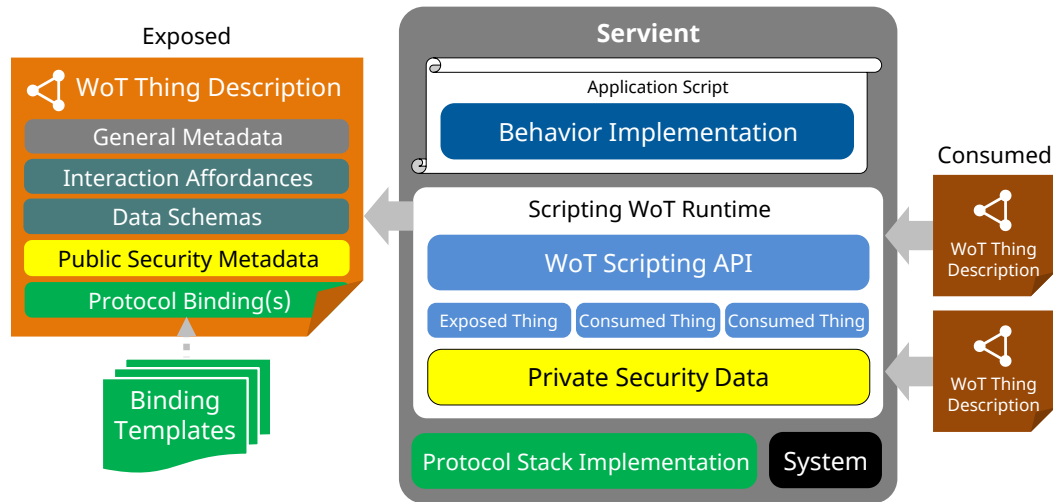


Figure 27 Implementation of a Servient using the WoT Scripting API

The role and functionality of each module shown in [Figure 27](#) is explained in the following sections.

8.1 Behavior Implementation §

The *behavior* defines the overall application logic of a Thing, which has several aspects:

It includes *autonomous behavior* of Things (e.g., sampling of sensors or control loops for actuators), the *handlers* for Interaction Affordances (i.e., the concrete actions taken when an affordance is activated), *Consumer behavior* (e.g., controlling a Thing or realizing mashups), and *Intermediary behavior* (e.g., simply proxying a Thing or composing virtual entities). The behavior implementation within a Servient defines which Things, Consumers, and Intermediaries are hosted on this component.

[Figure 27](#) depicts Servients that are implementing the optional WoT Scripting API building block, where portable application scripts written in JavaScript [[ECMAScript](#)] define the behavior. They are executed by a scripting runtime system that is part of the WoT Runtime (when providing the WoT Scripting API or any other script-based API). They are portable, as they are written against the common WoT Scripting API definitions, and hence can be executed by any Servient featuring this building block. This makes it possible to shift application logic between system components, for instance moving a Consumer from the cloud to an edge node to meet networking requirements, or to

move an Intermediary to the cloud to fulfill growing resource demands. Portable applications enable to 'install' additional behavior after the deployment of a Servient.

In principle, any programming language and API can be used in order to define the behavior of a Thing, as long as the Interaction Affordances are presented externally through a WoT Interface. The adaption between application-facing API and the protocol stack is handled by the WoT Runtime. See [§ 8.8.1 Native WoT API](#) for behavior implementation without the WoT Scripting API building block.

8.2 WoT Runtime §

Technically, the Thing abstraction and its Interaction Model is implemented in a runtime system. This WoT Runtime maintains the execution environment for the behavior implementation and is able to expose and/or consume Things, and hence must be able to fetch, process, serialize, and serve WoT Thing Descriptions.

Every WoT Runtime has an application-facing interface (i.e., an API) for the behavior implementation. The optional WoT Scripting API building block shown in [Figure 27](#) defines such an application-facing interface that follows the Thing abstraction and enables the deployment of behavior implementations during runtime through application scripts. See [§ 8.8.1 Native WoT API](#) for alternative APIs, which can also only be available during compile time. In general, application logic should be executed in isolated execution environments to prevent unauthorized access to the management aspects of the WoT Runtime, in particular the Private Security Data. In multi-tenant Servients, additional execution environment isolation is required for the different tenants.

A WoT Runtime needs to provide certain operations to manage the lifecycle of Things, or more precisely their software abstractions and descriptions. A lifecycle management (LCM) system may encapsulate those lifecycle operations within a Servient and use internal interfaces to realize the lifecycle management. The details of such operations vary among different implementations. The WoT Scripting API includes LCM functionality, and hence represents one possible implementation of such a system.

The WoT Runtime must interface with the protocol stack implementation of the Servient, as it decouples the behavior implementation from the details of the Protocol Bindings. The WoT Runtime usually also interfaces with the underlying system, for instance, to access local hardware such as attached sensors and actuators or to access system services such as storage. Both interfaces are implementation-specific, yet the WoT Runtime must provide the necessary adaption to the implemented Thing abstraction.

8.3 WoT Scripting API §

The WoT Scripting API building block defines an ECMAScript API that closely follows the WoT Thing Description specification [[WOT-THING-DESCRIPTION](#)]. It defines the interface between behavior implementations and a scripting-based WoT Runtime. Other, simpler APIs may be implemented on top of it, similar to, for instance, jQuery for the Web browser APIs.

See [[WOT-SCRIPTING-API](#)] for more details.

8.4 Exposed Thing and Consumed Thing Abstractions §

The WoT Runtime instantiates software representations of Things based on their TDs. These software representations provide the interface towards the behavior implementation.

The Exposed Thing abstraction represents a Thing hosted locally and accessible from the outside through the protocol stack implementation of the Servient. The behavior implementation can fully control Exposed Things by defining their metadata and Interaction Affordances, and providing their autonomous behavior.

The Consumed Thing abstraction represents a remotely hosted Thing for Consumers that needs to be accessed using a communication protocol. Consumed Things are proxy objects or stubs. The behavior implementation is restricted to reading their metadata and activating their Interaction Affordances as described in the corresponding TD. Consumed Things can also represent system features such as local hardware or devices behind proprietary or legacy communication protocols. In this case, the WoT Runtime must provide the necessary adaptation between system API and Consumed Thing. Furthermore, it must provide corresponding TDs and make them available to the behavior implementation, for instance, by extending whatever discovery mechanism is provided by the WoT Runtime through the application-facing API (e.g., the `discover()` method defined in the WoT Scripting API [[WOT-SCRIPTING-API](#)]).

When using the WoT Scripting API, Exposed Thing and Consumed Thing are JavaScript objects, which can be created, operated on, and destroyed by application scripts. However, access may be restricted through a security mechanism, for instance, in multi-tenant Servients.

8.5 Private Security Data §

Private security data, such as a secret key for interacting with the Thing, is also conceptually managed by the WoT Runtime, but is intentionally not made directly accessible to the application. In fact, in the most secure hardware implementations, such Private Security Data is stored in a separate, isolated

memory (e.g., on a secure processing element or TPM) and only an abstract set of operations (possibly even implemented by an isolated processor and software stack) is provided that limit the attack surface and prevent external disclosure of this data. [Private Security Data](#) is used transparently by the [Protocol Binding](#) to authorize and protect the integrity and confidentiality of interactions.

8.6 Protocol Stack Implementation §

The protocol stack of a [Servient](#) implements the [WoT Interface](#) of the [Exposed Things](#) and is used by [Consumers](#) to access the [WoT Interface](#) of remote [Things](#) (via [Consumed Things](#)). It produces the concrete protocol messages to interact over the network. [Servients](#) may implement multiple protocols, and hence support multiple [Protocol Bindings](#) to enable interaction with different [IoT Platforms](#).

In many cases, where standard protocols are used, generic protocol stacks can be used to produce the platform-specific messages (e.g., one for HTTP(S) dialects, one for CoAP(S) dialects, and one for MQTT solutions, etc.). In this case, the communication metadata from the [Thing Description](#) is used to select and configure the right stack (e.g., HTTP with the right header fields or CoAP with the right options). Parsers and serializers for the expected payload representation format (JSON, CBOR, XML, etc.) as identified by the media type [\[RFC2046\]](#) can also be shared across these generic protocol stacks.

See [\[WOT-BINDING-TEMPLATES\]](#) for details.

8.7 System API §

An implementation of a [WoT Runtime](#) may provide local hardware or system services to behavior implementations through the [Thing](#) abstraction, as if they were accessible over a communication protocol. In this case, the [WoT Runtime](#) should enable the behavior implementation to instantiate [Consumed Things](#) that internally interface with the system instead of the protocol stack. This can be done by listing such system Things, which are only available in the local [WoT Runtime](#), in the results of the discovery mechanism provided by the application-facing [WoT Runtime API](#).

A device may also be physically external to a [Servient](#), but connected via a proprietary protocol or a protocol not eligible as [WoT Interface](#) (see § 6.6 [Protocol Bindings](#)). In this case, the [WoT Runtime](#) may access legacy devices with such protocols (e.g., ECHONET Lite, BACnet, X10, I2C, SPI, etc.) through proprietary APIs, but may again choose to expose them to the behavior implementation via a [Thing](#) abstraction. A [Servient](#) can then act as gateway to the legacy devices. This should only be done if the legacy device cannot be described directly using a [WoT Thing Description](#).

The behavior implementation may also access local hardware or system services (e.g., storage) through a proprietary API or other means. This is, however, out of scope of the W3C WoT standardization, as it hinders portability.

8.8 Alternative Servient and WoT Implementations §

The WoT Scripting API building block is optional. Alternative Servient implementations are possible, where the WoT Runtime offers an alternative API for the application logic, which may be written in any programming language.

Furthermore, devices or services unaware of W3C WoT can still be consumed, when it is possible to provide a well-formed WoT Thing Description for them. In this case, the TD describes a WoT Interface of a Thing that has a black-box implementation.

8.8.1 Native WoT API §

There are various reasons why a developer may choose to implement a Servient without using the WoT Scripting API. This may be due to insufficient memory or computing resources, so the developer cannot use the required software stack or a fully-featured scripting engine. Alternatively, to support their use case (for example, a proprietary communications protocol) the developer may have to use specific functions or libraries only available through a particular programming environment or language.

In this case, a WoT Runtime can still be used, but with an equivalent abstraction and functionality exposed using an alternative application-facing interface instead of the WoT Scripting API. Except for the latter, all block descriptions in § 8. Abstract Servient Architecture are also valid for Figure 28.

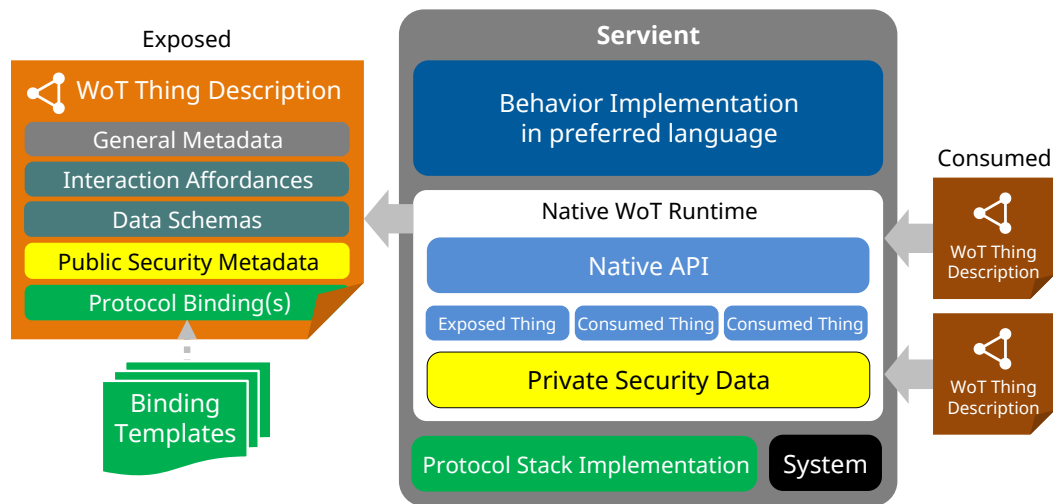


Figure 28 Implementation of a Servient Using a Native WoT API

8.8.2 Thing Description for Existing Devices §

It is also possible to integrate *existing* IoT devices or services into the W3C Web of Things and to use them as Things by creating a Thing Description for these devices or services. Such a TD can either be created manually or via a tool or service. For example, a TD could be generated by a service that provides automatic translation of metadata provided by another, ecosystem-dependent machine-readable format. This can only be done, however, if the target device is using protocols that can be described using a Protocol Binding. The requirements for this are given in § 6.6 Protocol Bindings. Much of the previous discussion also implies that a Thing provides its own Thing Description. While this is a useful pattern it is not mandatory. In particular, it may not be possible to modify existing devices to provide their own Thing Description directly. In this case the Thing Description will have to be provided separately using a service such as a directory or some other external and separate distribution mechanism.

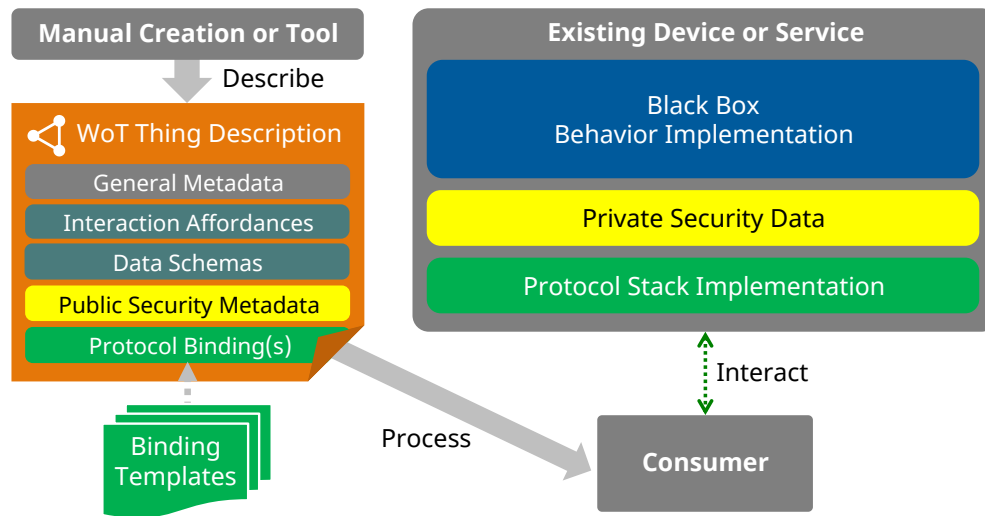


Figure 29 Integration of Existing IoT Devices into W3C WoT

9. Example WoT Deployments §

This section is non-normative.

This section provides various examples of how the Web of Things (WoT) abstract architecture may be instantiated when devices and services that implement the Thing and Consumer roles are connected together in various concrete topologies and deployment scenarios. These topologies and scenarios are not normative, but are permitted and supported by the WoT abstract architecture.

Before discussing specific topologies, we will first review the roles that Things and Consumers can play in a WoT network and the relationships they have with the Exposed Thing and Consumed Thing software abstractions. Exposed Thing and Consumed Thing are internally available to the behavior implementations of Servients in the roles of Things and Consumers, respectively.

9.1 Thing and Consumer Roles §

A Servient in the role of a Thing creates an Exposed Thing based on a Thing Description (TD). TDs are published and made available to other Servients that are in the roles of Consumers or Intermediaries. TDs may be published in various different ways: the TD might be registered with a management system such as a Thing Directory service, or a Thing may provide the requesters with a TD upon receiving a request for a TD. It is even possible to statically associate a TD with Thing in certain application scenarios.

A Servient in the role of a Consumer obtains the TD of a Thing using a discovery mechanism and creates a Consumed Thing based on the obtained TD. The concrete discovery mechanism depends on the individual deployment scenario: It could be provided by a management system such as a Thing Directory, a discovery protocol, through static assignment, etc.

However, it should be noted that TDs describing devices associated with an identifiable person may potentially be used to infer privacy-sensitive information. Constraints on the distribution of such TDs must therefore be incorporated into any concrete TD discovery mechanism. If possible, steps to limit the information exposed in a TD may also have to be taken, such as filtering out IDs or human-readable information if this is not strictly necessary for a particular use case. Privacy issues are discussed at a high level in § 10. [Security and Privacy Considerations](#) and a more detailed discussion is given in the [\[WOT-THING-DESCRIPTION\]](#) specification.

Internal system functions of a device, such as interacting with attached sensors and actuators, can also optionally be represented as Consumed Thing abstractions.

The functions supported by the Consumed Thing are provided to the Consumer's behavior implementation through a programming language interface. In the WoT Scripting API, Consumed Things are represented by objects. The behavior implementation (that is, the application logic) running in a Thing can engage through Interaction Affordances with Consumers by using the programming language interface provided by the Exposed Thing.

A Thing does not necessarily represent a physical device. Things can also represent a collection of devices, or virtual services running in a gateway or in the cloud. Likewise, a Consumer may represent an application or service running on a gateway or cloud. Consumers can also be implemented on edge devices. In Intermediaries, a single Servient performs both the roles of a Thing and a Consumer simultaneously which are sharing a single WoT Runtime.

9.2 Topology of WoT Systems and Deployment Scenarios §

Various topologies and deployment scenarios of WoT systems are discussed in this section. These are only example patterns and other interconnection topologies are also possible. The topologies described here are derived from the Web of Things use cases (§ 4. [Use Cases](#)) as well as the technical requirements extracted from them (§ 5. [Requirements](#)).

9.2.1 Consumer and Thing on the Same Network §

In the simplest interconnection topology, illustrated by [Figure 30](#), the Consumer and Thing are on the same network and can communicate directly with each other without any intermediaries. One use case

where this topology arises is when the Consumer is an orchestration service or some other IoT application running on a gateway and the Thing is a device interfacing to a sensor or an actuator. However, the client/server relationship could easily be reversed; the client could be a device in the Consumer role accessing a service running as a Thing on a gateway or in the cloud.



Figure 30 Consumer and Thing on the Same Network

If the Thing is in the cloud and the Consumer is on a local network (see [Figure 1](#) for an example in a Smart Home use case) the actual network topology may be more complex, for example requiring NAT traversal and disallowing certain forms of discovery. In such cases one of the more complex topologies discussed later may be more appropriate.

9.2.2 Consumer and Thing Connected via Intermediaries §

An Intermediary plays both Thing and Consumer roles on the network and supports both the Exposed Thing and Consumed Thing software abstractions within its WoT Runtime. Intermediaries can be used for proxying between devices and networks or for Digital Twins.

9.2.2.1 Intermediary Acting as a Proxy §

One simple application of an Intermediary is a proxy for a Thing. When the Intermediary acts as a proxy, it has interfaces with two separate networks or protocols. This may involve the implementation of additional security mechanisms such as providing TLS endpoints. Generally proxies do not modify the set of interactions, so the TD exposed by the Intermediary will have the same interactions as the consumed TD, however the connection metadata is modified.

To implement this proxy pattern, the Intermediary obtains a TD of a Thing and creates a Consumed Thing. It creates a proxy object of the Thing as a software implementation that has the same Interaction Affordances. It then creates a TD for the proxy object with a new identifier and possibly with new communications metadata (Protocol Bindings) and/or new Public Security Metadata. Finally, an Exposed Thing is created based on this TD, and the Intermediary notifies other Consumers or Intermediaries of the TD via an appropriate publication mechanism.




Figure 31 Consumer and Thing Connect via an Intermediary Acting as a Proxy

9.2.2.2 Intermediary Acting as a Digital Twin §

More complex Intermediaries may be known as Digital Twins. A Digital Twin may or may not modify the protocols or translate between networks, but they provide additional services, such as state caching, deferred updates, or even predictive simulation of the behavior of the target device. For example, if an IoT device has limited power, it may choose to wake up relatively infrequently, synchronize with a digital twin, and then immediately go to sleep again. In this case, typically the Digital Twins runs on a less power-constrained device (such as in the cloud or on a gateway) and is able to respond to interactions on the constrained device's behalf. Requests for the current state of properties may also be satisfied by the Digital Twins using cached state. Requests that arrive when the target IoT device is sleeping may be queued and sent to it when it wakes up. To implement this pattern, the Intermediary, i.e., the digital twin needs to know when the device is awake. The device implementation as a Thing may need to include a notification mechanism for that. This could be implemented using a separate Consumer/Thing pair, or by using Event interactions for this purpose.

9.2.3 Devices in a Local Network Controlled from a Cloud Service §

In Smart Home use cases, devices (sensors and home appliances) connected to a home network are often monitored and, in some cases, also controlled by cloud services. There is usually a NAT device between the home network to which the devices are connected and the cloud. The NAT device translates IP addresses as well as often providing firewall services, which block connections selectively. The local devices and cloud services can only communicate with each other if the communication can successfully traverse the gateway.

A typical structure, adopted in ITU-T Recommendation Y.4409/Y.2070 [Y.4409-Y.2070], is shown in [Figure 32](#). In this structure there is both a local and a remote Intermediary. The local Intermediary aggregates the Interaction Affordances from multiple Thing into a (set of) Exposed Things, which can all be mapped onto a common protocol (for example, HTTP, with all interactions mapped to a single URL namespace with a common base server and using a single port). This provides the remote Intermediary with a simple way to access all the Things behind the NAT device, assuming the local Intermediary has used a converged protocol that can traverse the NAT device and has some way to expose this service to the Internet (STUN, TURN, DyDNS, etc.). In addition, the local Intermediary

can function as a Thing proxy, so even when the connected Things each use a different protocol (HTTP, MQTT, CoAP, etc.) and/or a different set of ecosystem conventions, the Exposed Thing can converge them into a single protocol so that Consumers do not need to be aware of the various protocols the Things use.

In [Figure 32](#), there are two clients connected to the remote Intermediary, which has aggregated the services that reside behind the NAT border and may provide additional protocol translation or security services. In particular, the local Intermediary may be on a network with limited capacity and making that service directly available to all users may not be feasible. In this case access to the local Intermediary is *only* provided to the remote Intermediary. The remote Intermediary then implements a more general access control mechanism and may also perform caching or throttling to protect the consumer from excess traffic. Those consumers also will use a single protocol suitable for the open Internet (e.g., HTTPS) to communicate with the Intermediary, which makes the development of clients much simpler.

In this topology there is NAT and firewall functionality between the consumers and things, but the local and remote Intermediaries work together to tunnel all communications through the firewall, so the consumers and things need to know nothing about the firewall. The paired Intermediaries also protect the home devices by providing access control and traffic management.




Figure 32 Cloud Applications Implemented as Consumers Connected to Local Devices implemented as Things via Paired Intermediaries

In more difficult cases the NAT and firewall traversal may not work exactly as shown. In particular, an ISP may not support publicly accessible addresses, or STUN/TURN and/or DyDNS may not be supported or available. In this case the Intermediaries may alternatively reverse the client/server roles between them to set up an initial connection (with the local Intermediary first connecting to the remote Intermediary in the cloud), then the pair of Intermediaries may establish a tunnel (using for example, a Secure WebSocket, which uses TLS to protect the connection). The tunnel can then be used to encode all communications between the Intermediaries using a custom protocol. In this case the initial

connection can still be made over HTTPS using standard ports, and from the local Intermediary to the remote Intermediary identically to a normal browser/web server interaction. This should be able to traverse most home firewalls, and since the connection is outgoing, network address translation will not cause any problems. However, even though a custom tunneling protocol is needed, the remote Intermediary can still translate this custom protocol back into standard external protocols. The connected Consumers and Things do not need to know about it. It is also possible to extend this example to use cases where both Things and Consumers can connect on either side of the NAT boundary. This however also requires a bidirectional tunnel to be established between the two Intermediaries.

9.2.4 Discovery Using a Thing Directory §

Once local devices (and possibly services) can be monitored or controlled by services on cloud, a variety of additional services can be built on top. For example, a cloud application could change a device's operating condition based on an analysis of collected data.

However when the remote Intermediary is a part of a cloud platform servicing client applications, the clients need to be able to find device information by, for example, accessing a directory of connected devices. For simplicity in the figure below we have assumed all local devices are implemented as Things and all cloud applications as Consumers. To make the metadata of local devices implemented as Things available to the cloud applications, their metadata can be registered with a Thing Directory service. This metadata is specifically the TDs of the local devices modified to reflect the Public Security Metadata and communication metadata (Protocol Bindings) provided by the remote Intermediary. A client application then can obtain the metadata it needs to communicate with local devices to achieve its functionality by querying the Thing Directory.




Figure 33 Cloud Service with Thing Directory

In more complex situations, not shown in the figure, there may also be cloud services that act as Things. These can also register themselves with a Thing Directory. Since a Thing Directory is a Web service, it should be visible to the local devices through the NAT or firewall device and its interface

can even be provided with its own TD. Local devices acting as Consumers can then discover the Things in the cloud via a Thing Directory and connect to the Things directly or via the local Intermediary if, for instance, protocol translation is needed.

9.2.5 Service-to-Service Connections Across Multiple Domains §

Multiple cloud eco-systems each based on different IoT platforms can work together to make a larger, system-of-systems eco-system. Building on the previously discussed structure of a cloud application eco-system, the figure below shows two eco-systems connected to each other to make a system-of-systems. Consider the case in which a client in one eco-system (i.e., Consumer A below) needs to use a server in another eco-system (i.e., Thing B below). There is more than one mechanism to achieve this cross eco-systems application-device integration. Below, two mechanisms are explained, each using a figure, to show how this can be achieved.

9.2.5.1 Connection Through Thing Directory Synchronization §

In [Figure 34](#), two Thing Directories synchronize information, which makes it possible for Consumer A to obtain the information of Thing B through Thing Directory A. As described in previous sections, remote Intermediary B maintains a shadow implementation of Thing B. By obtaining the TD of this shadow device, Consumer A is able to use Thing B through the remote Intermediary B.

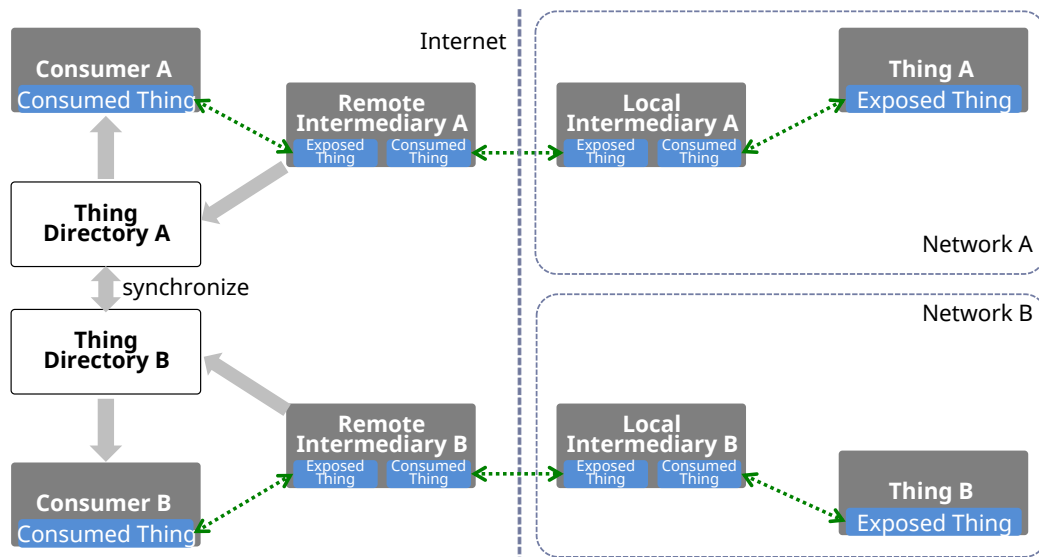


Figure 34 Multiple Cloud Connections Through Thing Directory Synchronization

In [Figure 35](#), two remote [Intermediaries](#) synchronize device information. When a shadow of [Thing B](#) is created in remote [Intermediary B](#), the shadow's TD is simultaneously synchronized into remote [Intermediary A](#). Remote [Intermediary A](#) in turn creates its own shadow of [Thing B](#), and registers the TD with [Thing Directory A](#). With this mechanism, synchronization between [Thing Directories](#) is not necessary.

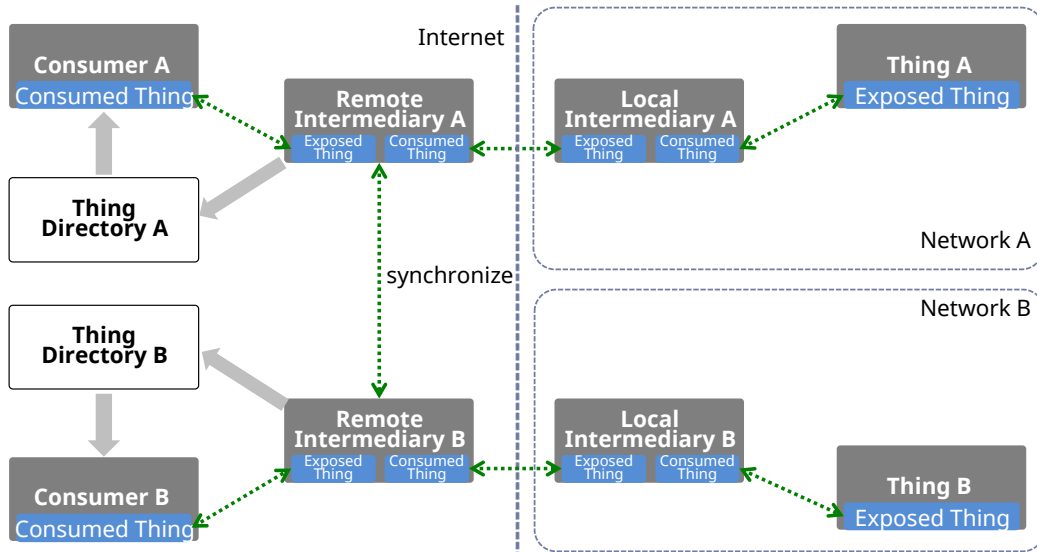


Figure 35 Multiple Cloud Connections Through Intermediary Synchronization

10. Security and Privacy Considerations §

This section is non-normative.

Security and privacy are a cross-cutting issues that need to be considered in all [WoT building blocks](#) and WoT implementations. This chapter summarizes some general issues and guidelines to help preserve the security and privacy of concrete WoT implementations. However, these are only general guidelines and an abstract architecture such as described in this document cannot, itself, guarantee security and privacy. Instead the details of a concrete implementation need to be considered. For a more detailed and complete analysis of security and privacy issues, see the *WoT Security and Privacy Guidelines* specification [\[WOT-SECURITY\]](#).

Overall, the goal of the WoT is to describe the existing access mechanisms and properties of IoT devices and services, including security. In general, W3C WoT is designed to describe what exists rather than to prescribe what to implement. A description of an existing system should accurately

describe that system, even if it has less than ideal security behavior. A clear understanding of the security vulnerabilities of a system supports security mitigations—although of course such data need not be distributed to those who might maliciously exploit it.

However, especially for new systems, the WoT architecture should *enable* the use of best practices in security and privacy. In general, the WoT security architecture must support the goals and mechanisms of the IoT protocols and systems it connects to. These systems vary in their security requirements and risk tolerance, so security mechanisms will also vary based on these factors.

Security and privacy are especially important in the IoT domain since IoT devices need to operate autonomously and, in many cases, have access to both personal data and/or can be in control of safety-critical systems. IoT devices are subject to different and in some cases higher risks than IT systems. It is also important to protect IoT systems so that they cannot be used to launch attacks on other computer systems.

In general, security and privacy cannot be guaranteed. It is not possible for the WoT to turn an insecure system into a secure one. However, the WoT architecture needs to do no harm: it should support security and privacy at least as well as the systems it describes and supports.

10.1 WoT Thing Description Risks §

The metadata contained in a WoT Thing Description (TD) is potentially sensitive. As a best practice, TDs should be used together with integrity protection mechanisms and access control policies, and should be provided only to authorized users.

Please refer to the Security and Privacy Consideration section of the WoT Thing Description specification for additional details and discussion of these points.

10.1.1 Thing Description Private Security Data Risk §

TDs are designed to carry only Public Security Metadata. Producers of TDs must ensure that no Private Security Data is included in TDs. There should be a strict separation of Public Security Metadata and Private Security Data. A TD should contain only Public Security Metadata, letting Consumers know what they need to do to access a system if and only if they have authorization. Authorization in turn should be based on separately managed private information.

The built-in TD security schemes defined in the TD specification intentionally do not support the encoding of Private Security Data. However, there is a risk that other fields such as human-readable

descriptions might be misused to encode this information, or new security schemes might be defined and deployed via the extension mechanism that encode such information.

Mitigation:

Creators of TDs and extensions meant to be used in TDs must ensure that only Public Security Metadata is ever stored in TDs.

10.1.2 Thing Description Personally Identifiable Information Risk §

Thing descriptions can potentially contain Personally Identifiable Information of various types. Even if it is not explicit, a TD and its association with an identifiable person can be used to infer information about that person. For example, the association of fingerprintable TDs exposed by mobile devices whose location can be determined can be a tracking risk. Even if a particular device instance cannot be identified, the type of device represented by a TD, when associated with a person, may constitute personal information. For example, a medical device may be used to infer that the user has a medical condition.

Generally, Personally Identifiable Information in a TD should be limited as much as possible. In some cases, however, it cannot be avoided. The potential presence of both direct and inferencable PII in a TD means that TD should be treated like other forms of PII. They should be stored and transmitted in a secure fashion, should only be provided to authorized users, should only be cached for limited times, should be deleted upon request, should only be used for the purpose for which they were provided with user consent, and they should otherwise satisfy all requirements (including any legal requirements) for the use of PII.

Mitigation:

Storage of PII in TDs should be minimized as much as possible. Even without explicit PII in TDs, a tracking and identification privacy risk may exist. To minimize this risk, TDs should generally be treated as if they contained PII and subject to the same management policies as other PII. They should only be provided to authorized Consumers. Information unnecessary for a specific use case should not be exposed in TDs whenever possible. For example, explicit type and instance identifying information in TDs should also not be included if it is not needed by the use case. Even if required by the use case, to minimize tracking risks, distributed and limited-scope identifiers should be used whenever possible rather than globally unique identifiers. Other forms of information, such as human-readable descriptions, may also be omitted in some use cases to reduce fingerprinting risks.

10.1.3 Thing Description Communication Metadata Risk §

The WoT Binding Templates must correctly support the security mechanisms employed by the underlying IoT Platform for that platform to be considered eligible for use with WoT. Due to the automation of network interactions necessary to deploy IoT at scale, operators need to ensure that Things are exposed and consumed in a way that is compliant with their security policies.

Mitigation:

Whenever possible, TD creators should use the vetted communication metadata provided in the WoT Binding Templates. When generating TDs for an IoT ecosystem not covered by the WoT Binding Templates, ensure that all the security requirements of the IoT Platform are satisfied.

10.2 WoT Scripting API Security and Privacy Risks §

The WoT Runtime implementation and the WoT Scripting API should have mechanisms to prevent malicious access to the system and isolate scripts in multi-tenant Servients. More specifically the WoT Runtime implementation when used with the WoT Scripting API should consider the following security and privacy risks and implement the recommended mitigations.

10.2.1 Cross-Script Security and Privacy Risk §

In basic WoT setups, all scripts running inside the WoT Runtime are considered trusted, distributed by the manufacturer, and therefore there is no strong need to perform strict isolation between each running script instance. However, depending on device capabilities, deployment use case scenarios, and risk level it might be desirable to do so. For example, if one script handles sensitive privacy-related PII data and is well-audited, it might be desirable to separate it from the rest of the script instances to minimize the risk of data exposure in case some other script inside the same system gets compromised during runtime. Another example is mutual co-existence of different tenants on a single WoT device. In this case each WoT runtime instance will be hosting a different tenant, and isolation between them is required.

Mitigation:

The WoT Runtime should perform isolation of script instances and their data in cases when scripts handle privacy-related or other critical security data. Similarly, the WoT Runtime implementation should perform isolation of WoT Runtime instances and their data if a WoT device has more than one tenant. Such isolation can be performed within the WoT Runtime using platform security mechanisms available on the device. For more information see Sections "WoT Servient Single-Tenant" and "WoT Servient Multi-Tenant" of the *WoT Security and Privacy Guidelines* specification [[WOT-SECURITY](#)].

10.2.2 Physical Device Direct Access Security and Privacy Risk §

In case a script is compromised or malfunctions the underlying physical device (and potentially surrounded environment) can be damaged if a script can use directly exposed native device interfaces. If such interfaces lack safety checks on their inputs, they might bring the underlying physical device (or environment) to an unsafe state.

Mitigation:

The WoT Runtime should avoid directly exposing the native device interfaces to the script developers. Instead a WoT Runtime implementation should provide a hardware abstraction layer for accessing the native device interfaces. Such hardware abstraction layer should refuse to execute commands that might put the device (or environment) to an unsafe state. Additionally, in order to reduce the damage to a physical WoT device in cases a script gets compromised, it is important to minimize the number of interfaces that are exposed or accessible to a particular script based on its functionality.

10.3 WoT Runtime Security and Privacy Risks §

10.3.1 Provisioning and Update Security Risk §

If the WoT Runtime implementation supports post-manufacturing provisioning or updates of itself, scripts, or any related data (including security credentials), it can be a major attack vector. An attacker can try to modify any above described element during the update or provisioning process or simply provision attacker's code and data directly.

Mitigation:

Post-manufacturing provisioning or update of scripts, the WoT Runtime itself or any related data should be done in a secure fashion. A set of recommendations for secure update and post-manufacturing provisioning can be found in the *WoT Security and Privacy Guidelines* specification [[WOT-SECURITY](#)].

10.3.2 Security Credentials Storage Security and Privacy Risk §

Typically the WoT Runtime needs to store the security credentials that are provisioned to a WoT device to operate in a network. If an attacker can compromise the confidentiality or integrity of these credentials, then it can obtain access to assets, impersonate other WoT Things, devices, or services, or launch Denial-Of-Service (DoS) attacks.

Mitigation:

The WoT Runtime should securely store any provisioned security credentials, guaranteeing their integrity and confidentiality. In case there are more than one tenant on a single WoT-enabled device, a WoT Runtime implementation should guarantee isolation of each tenant's provisioned security credentials. Additionally, in order to minimize a risk that provisioned security credentials get compromised, the WoT Runtime implementation should not expose any API for scripts to query the provisioned security credentials. Such credentials (or even better, abstract operations that use them but do not expose them) should only be accessible to the Protocol Binding implementation that uses them.

A. Recent Specification Changes §

Changes from Proposed Recommendation §

- No normative changes, minor editorial fixes and stable external references.

Changes from First Candidate Recommendation §

- Make Terminology section informative.
- Reorganized abstract and introduction.
- Expanded privacy discussion and mitigations.
- Adapted figures to text changes.
- Updated and extended definitions.
- Updated references:
 - Normative references added: RFC2046
 - Normative references removed or moved to informative section: IANA-RELATIONS, MQTT, RFC4395, RFC6838, RFC7049, RFC7231, RFC7252
- Adjusted figure colors based on accessibility feedback.
- Minor editorial fixes such as typos and capitalization.

Changes from First Public Working Draft §

- Revised and expanded use cases.

- Reorganized requirements.
- Definition of abstract architecture.
- Revised and clarified terminology.
- Additions to implementation and deployments.
- Addition of security and privacy considerations.

B. Acknowledgments §

Special thanks to Michael McCool, Takuki Kamiya, Kazuyuki Ashimura, Sebastian Käbisch, Zoltan Kis, Elena Reshetova, Klaus Hartke, Ari Keränen, Kazuaki Nimura, and Philippe Le Hegaret for their contributions to this document.

Many thanks to the W³C staff and all other active Participants of the W³C Web of Things Interest Group (WoT IG) and Working Group (WoT WG) for their support, technical input and suggestions that led to improvements to this document.

The WoT WG also would like to appreciate the pioneering efforts regarding the concept of "Web of Things" that started as an academic initiative in the form of publications such as [[WOT-PIONEERS-1](#)] [[WOT-PIONEERS-2](#)] [[WOT-PIONEERS-3](#)] [[WOT-PIONEERS-4](#)] and, starting in 2010, a yearly [International Workshop on the Web of Things](#).

Finally, special thanks to Joerg Heuer for leading the WoT IG for 2 years from its inception and guiding the group to come up with the concept of WoT building blocks including the Thing Description.

C. References §

C.1 Normative references §

[RFC2046]

Multipurpose Internet Mail Extensions (MIME) Part Two: Media Types. N. Freed; N. Borenstein. IETF. November 1996. Draft Standard. URL: <https://tools.ietf.org/html/rfc2046>

[RFC2119]

Key words for use in RFCs to Indicate Requirement Levels. S. Bradner. IETF. March 1997. Best Current Practice. URL: <https://tools.ietf.org/html/rfc2119>

[RFC3986]

Uniform Resource Identifier (URI): Generic Syntax. T. Berners-Lee; R. Fielding; L. Masinter. IETF. January 2005. Internet Standard. URL: <https://tools.ietf.org/html/rfc3986>

[RFC3987]

Internationalized Resource Identifiers (IRIs). M. Duerst; M. Suignard. IETF. January 2005. Proposed Standard. URL: <https://tools.ietf.org/html/rfc3987>

[RFC5234]

Augmented BNF for Syntax Specifications: ABNF. D. Crocker, Ed.; P. Overell. IETF. January 2008. Internet Standard. URL: <https://tools.ietf.org/html/rfc5234>

[RFC8174]

Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words. B. Leiba. IETF. May 2017. Best Current Practice. URL: <https://tools.ietf.org/html/rfc8174>

[RFC8259]

The JavaScript Object Notation (JSON) Data Interchange Format. T. Bray, Ed.. IETF. December 2017. Internet Standard. URL: <https://tools.ietf.org/html/rfc8259>

[RFC8288]

Web Linking. M. Nottingham. IETF. October 2017. Proposed Standard. URL: <https://httpwg.org/specs/rfc8288.html>

C.2 Informative references §

[CoRAL]

The Constrained RESTful Application Language (CoRAL). Klaus Hartke. IETF. March 2019. Internet-Draft. URL: <https://tools.ietf.org/html/draft-hartke-t2trg-coral>

[CoRE-RD]

CoRE Resource Directory. M. Koster; C. Bormann; P. van der Stok; C. Amsuess. IETF. 13 June 2019. Internet-Draft. URL: <https://tools.ietf.org/html/draft-ietf-core-resource-directory-21>

[ECMAScript]

ECMAScript Language Specification. Ecma International. URL: <https://tc39.es/ecma262/>

[HCI]

The Encyclopedia of Human-Computer Interaction, 2nd Ed. Interaction Design Foundation. 2013. URL: <https://www.interaction-design.org/literature/book/the-encyclopedia-of-human-computer-interaction-2nd-ed>

[HTML]

HTML Standard. Anne van Kesteren; Domenic Denicola; Ian Hickson; Philip Jägenstedt; Simon Pieters. WHATWG. Living Standard. URL: <https://html.spec.whatwg.org/multipage/>

[IANA-RELATIONS]

Link Relations. IANA. URL: <https://www.iana.org/assignments/link-relations/>

[IANA-URI-SCHEMES]

Uniform Resource Identifier (URI) Schemes. IANA. URL: <https://www.iana.org/assignments/uri-schemes/uri-schemes.xhtml>

[IEC-FOTF]

Factory of the future. IEC. October 2015. URL: <https://www.iec.ch/whitepaper/pdf/iecWP-futurefactory-LR-en.pdf>

[IOT-SCHEMA-ORG]

Schema Extensions for IoT Community Group. URL: <https://www.w3.org/community/iotschema/>

[ISO-IEC-2382]

Information technology — Vocabulary. ISO. 2015. URL: <https://www.iso.org/obp/ui/#iso:std:iso-iec:2382:ed-1:v1:en>

[ISO-IEC-27000]

Information technology — Security techniques — Information security management systems — Overview and vocabulary. ISO. 2018. URL: <https://www.iso.org/obp/ui/#iso:std:iso-iec:27000:ed-5:v1:en>

[ISO-IEC-29100]

Information technology — Security techniques — Privacy framework. ISO. 2011. URL: <https://www.iso.org/obp/ui/#iso:std:iso-iec:29100:ed-1:v1:en>

[JSON-LD11]

JSON-LD 1.1. Gregg Kellogg; Pierre-Antoine Champin; Dave Longley. W3C. 5 March 2020. W3C Candidate Recommendation. URL: <https://www.w3.org/TR/json-ld11/>

[LINKED-DATA]

Linked Data Design Issues. Tim Berners-Lee. W3C. 27 July 2006. W3C-Internal Document. URL: <https://www.w3.org/DesignIssues/LinkedData.html>

[LWM2M]

Lightweight Machine to Machine Technical Specification: Core. OMA SpecWorks. August 2018. Approved Version: 1.1. URL: http://openmobilealliance.org/release/LightweightM2M/V1_1-20180710-A/OMA-TS-LightweightM2M_Core-V1_1-20180710-A.pdf

[MQTT]

MQTT Version 3.1.1 Plus Errata 01. Andrew Banks; Rahul Gupta. OASIS Standard. December 2015. URL: <http://docs.oasis-open.org/mqtt/mqtt/v3.1.1/mqtt-v3.1.1.html>

[NORMAN]

The Psychology of Everyday Things. Donald A. Norman. Basic Books. 1988.

[OCF]

OCF Core Specification. Open Connectivity Foundation. April 2019. Version 2.0.2. URL: <https://openconnectivity.org/developer/specifications>

[REST]

REST: Architectural Styles and the Design of Network-based Software Architectures. Roy Thomas Fielding. University of California, Irvine. 2000. PhD thesis. URL: https://www.ics.uci.edu/~fielding/pubs/dissertation/fielding_dissertation.pdf

[RFC4301]

Security Architecture for the Internet Protocol. S. Kent; K. Seo. IETF. December 2005. Proposed Standard. URL: <https://tools.ietf.org/html/rfc4301>

[RFC6202]

Known Issues and Best Practices for the Use of Long Polling and Streaming in Bidirectional HTTP. S. Loreto; P. Saint-Andre; S. Salsano; G. Wilkins. IETF. April 2011. Informational. URL: <https://tools.ietf.org/html/rfc6202>

[RFC6347]

Datagram Transport Layer Security Version 1.2. E. Rescorla; N. Modadugu. IETF. January 2012. Proposed Standard. URL: <https://tools.ietf.org/html/rfc6347>

[RFC6690]

Constrained RESTful Environments (CoRE) Link Format. Z. Shelby. IETF. August 2012. Proposed Standard. URL: <https://tools.ietf.org/html/rfc6690>

[RFC6749]

The OAuth 2.0 Authorization Framework. D. Hardt, Ed.. IETF. October 2012. Proposed Standard. URL: <https://tools.ietf.org/html/rfc6749>

[RFC7049]

Concise Binary Object Representation (CBOR). C. Bormann; P. Hoffman. IETF. October 2013. Proposed Standard. URL: <https://tools.ietf.org/html/rfc7049>

[RFC7231]

Hypertext Transfer Protocol (HTTP/1.1): Semantics and Content. R. Fielding, Ed.; J. Reschke, Ed.. IETF. June 2014. Proposed Standard. URL: <https://httpwg.org/specs/rfc7231.html>

[RFC7252]

The Constrained Application Protocol (CoAP). Z. Shelby; K. Hartke; C. Bormann. IETF. June 2014. Proposed Standard. URL: <https://tools.ietf.org/html/rfc7252>

[RFC7641]

Observing Resources in the Constrained Application Protocol (CoAP). K. Hartke. IETF. September 2015. Proposed Standard. URL: <https://tools.ietf.org/html/rfc7641>

[RFC7744]

Use Cases for Authentication and Authorization in Constrained Environments. L. Seitz, Ed.; S. Gerdes, Ed.; G. Selander; M. Mani; S. Kumar. IETF. January 2016. Informational. URL: <https://tools.ietf.org/html/rfc7744>

[RFC8446]

The Transport Layer Security (TLS) Protocol Version 1.3. E. Rescorla. IETF. August 2018. Proposed Standard. URL: <https://tools.ietf.org/html/rfc8446>

[SAREF]

Smart Appliances REference (SAREF) ontology. ETSI. November 2015. URL: <https://sites.google.com/site/smartappliancesproject/ontologies/reference-ontology>

[VOCAB-SSN]

Semantic Sensor Network Ontology. Armin Haller; Krzysztof Janowicz; Simon Cox; Danh Le Phuoc; Kerry Taylor; Maxime Lefrançois. W3C. 19 October 2017. W3C Recommendation. URL: <https://www.w3.org/TR/vocab-ssn/>

[WOT-BINDING-TEMPLATES]

Web of Things (WoT) Binding Templates. Michael Koster; Ege Korkan. W3C. 30 January 2020. W3C Note. URL: <https://www.w3.org/TR/wot-binding-templates/>

[WOT-PIONEERS-1]

Mobile Service Interaction with the Web of Things. E. Rukzio, M. Paolucci; M. Wagner, H. Berndt; J. Hamard; A. Schmidt. Proceedings of 13th International Conference on Telecommunications (ICT 2006), Funchal, Madeira island, Portugal. May 2006. URL: <https://pdfs.semanticscholar.org/3ee3/a2e8ce93fbf9ba14ad54e12adaeb1f3ca392.pdf>

[WOT-PIONEERS-2]

Putting Things to REST. Erik Wilde. UCB iSchool Report 2007-015, UC Berkeley, Berkeley, CA, USA. November 2007. URL: <http://dret.net/netdret/docs/wilde-irep07-015-restful-things.pdf>

[WOT-PIONEERS-3]

Poster Abstract: Dyser – Towards a Real-Time Search Engine for the Web of Things. Benedikt Ostermaier; B. Maryam Elahi; Kay Römer; Michael Fahrmaier; Wolfgang Kellerer. Proceedings of ACM SenSys 2008, Raleigh, NC, USA. November 2008. URL: <https://www.vs.inf.ethz.ch/publ/papers/ostermair-poster-2008.pdf>

[WOT-PIONEERS-4]

A Resource Oriented Architecture for the Web of Things. Dominique Guinard; Vlad Trifa; Erik Wilde. Proceedings of Internet of Things 2010 International Conference (IoT 2010). Tokyo, Japan. November 2010. URL: <https://ieeexplore.ieee.org/abstract/document/5678452>

[WOT-SCRIPTING-API]

Web of Things (WoT) Scripting API. Zoltan Kis; Daniel Peintner; Johannes Hund; Kazuaki Nimura. W3C. 28 October 2019. W3C Working Draft. URL: <https://www.w3.org/TR/wot-scripting-api/>

[WOT-SECURITY]

Web of Things (WoT) Security and Privacy Guidelines. Elena Reshetova; Michael McCool. W3C. 6 November 2019. W3C Note. URL: <https://www.w3.org/TR/wot-security/>

[WOT-THING-DESCRIPTION]

Web of Things (WoT) Thing Description. Sebastian Käbisch; Takuki Kamiya; Michael McCool; Victor Charpenay; Matthias Kovatsch. W3C. 9 April 2020. W3C Recommendation. URL: <https://www.w3.org/TR/wot-thing-description/>

[Y.4409-Y.2070]

ITU-T Rec. Y.4409/Y.2070 (01/2015) Requirements and architecture of the home energy management system and home network services. ITU-T. January 2015. Recommendation. URL: <https://www.itu.int/rec/T-REC-Y.2070-201501-I>

